Quality in the Agent Skills Ecosystem: A Structural, Behavioral, and Cross-Contamination Analysis of 673 Skills

Dachary Carey

February 2026

Abstract

We present a systematic quality and content analysis of the Agent Skills ecosystem (modular instruction sets that extend AI coding agents), evaluating 673 skills from 41 source repositories. Complementing recent security-focused work that found 26.1% of skills contain security vulnerabilities (Liu et al. 2026), we focus on structural compliance, content quality, and cross-contamination risk. Our structural audit finds that 22.0% of skills fail validation, with company-published skills (79.2% pass rate) performing worse than community collections (94.0%). A token budget analysis reveals that 52% of all tokens across the ecosystem are nonstandard files wasting context window space. We identify 10 skills with high cross-contamination risk and 66 with "hidden contamination" visible only in reference files. LLM-as-judge scoring across all 673 skills reveals a two-factor quality structure: five craft dimensions intercorrelate strongly while novelty (information beyond training data) is largely independent, making it the key quality differentiator. An exploratory behavioral evaluation of 19 representative skills finds no correlation between structural contamination scores and measured behavioral degradation ($r = 0.077$, $n = 19$), and identifies six content-specific interference mechanisms through case studies: template propagation, textual frame leakage, token budget competition, API hallucination, cross-language code bleed, and architectural pattern bleed. These case studies suggest that the drivers of degradation are content-specific rather than language-mixing artifacts, and that realistic agentic context may substantially attenuate interference. We propose quality criteria for skill authors and recommendations for specification maintainers.

Contents

1 Introduction

The emergence of AI coding agents (tools like Claude Code, GitHub Copilot, and Cursor) has shifted software development toward human-AI collaboration. These agents operate in agentic loops: they read code,

plan modifications, write implementations, and run tests, with the human developer providing guidance and oversight.

Agent Skills ("Agent Skills Specification" 2025) extend this model by allowing developers and organizations to package domain expertise as modular instruction sets. A skill is a structured directory (primarily a `SKILL.md` markdown file plus optional reference materials) that an agent loads into its context window to gain specialized knowledge. Skills cover domains from specific tools (Playwright, Docker) to development methodologies (test-driven development, code review).

The Agent Skills ecosystem has grown rapidly. As of February 2026, the agentskills.io specification is supported by over 27 agent platforms including Claude Code, GitHub Copilot, Cursor, Windsurf, OpenAI Codex, and Gemini CLI. Companies from Microsoft and OpenAI to Stripe and Cloudflare publish official skills for their platforms. Community contributors have built skills for security analysis (Trail of Bits), scientific computing (K-Dense), and development workflows (Superpowers). We catalog 673 skills from 41 source repositories in our primary analysis, with an additional 800+ skills identified across the broader ecosystem.

However, this growth has outpaced quality assurance. Recent work has begun to address the security dimension: Liu et al. (Liu et al. 2026) analyzed 31,132 skills from marketplace aggregators and found that 26.1% contain security vulnerabilities spanning prompt injection, data exfiltration, privilege escalation, and supply chain risks. No prior work has systematically evaluated skill *quality*: the structural compliance, content effectiveness, and cross-contamination characteristics that determine whether skills actually improve agent output. We identified three concerns:

1. **Structural compliance**: Many skills deviate from the specification, potentially confusing agents about which files to load and how to use them.
2. **Content quality**: Skills vary enormously in how clearly and specifically they instruct agents, from highly structured step-by-step guides to vague advisory text.
3. **Cross-contamination risk**: Skills for multi-interface tools (databases, cloud services, container runtimes) may include examples in one language that interfere with the agent's code generation in another language, a concern supported by research on Programming Language Confusion (Moumoula et al. 2026), copy bias in in-context learning (Ali et al. 2024), and the demonstrated susceptibility of LLMs to irrelevant context (Shi et al. 2023).

This paper presents a systematic analysis across all three dimensions, using automated validation, content metrics, cross-contamination detection, and an exploratory behavioral evaluation that investigates what mechanisms drive degradation and whether structural risk metrics capture them. Our key contributions are:

- A comprehensive structural and quality audit of the Agent Skills ecosystem using the `skill-validator` tool (Carey 2026), covering 673 skills from 41 repositories, with a two-pass validation approach that separates deterministic structural checks from environment-dependent link validation. This complements the security-focused analysis of Liu et al. (Liu et al. 2026), which examined a larger sample (31,132 skills) for vulnerability patterns
- A token budget composition analysis revealing that 52% of all tokens across the ecosystem are non-standard files wasting context window space
- Content quality metrics (information density, instruction specificity) and LLM-as-judge scoring applied at ecosystem scale, revealing a two-factor quality structure where novelty is largely independent of craft dimensions
- Identification and taxonomy of cross-contamination risk in multi-interface skills, including the discovery of "hidden contamination" in reference files (66 skills with clean instruction files but contaminated references)

- An exploratory behavioral evaluation of 19 representative skills that identifies six distinct content interference mechanisms through case studies, finds no correlation between structural contamination scores and behavioral degradation (r = 0.077, n = 19), and provides preliminary evidence that realistic agentic context substantially attenuates measured degradation
- An ecosystem survey cataloging 1,400+ skills across 120+ repositories, contextualizing our findings as industry-wide
- Concrete recommendations for skill authors and specification maintainers

2 Methodology

2.1 Dataset

We collected 673 skills from 41 source repositories, organized into eight analytical categories:

Category	Skills	Repos	Description
Anthropic	16	1	Official skills from the spec authors
Company	288	22	Skills published by companies for their own APIs/products
Community collections	167	3	Multi-skill community repositories
Community individual	9	6	Single-skill community repositories
Trail of Bits	52	1	Security-focused vulnerability analysis skills
K-Dense (Superpowers)	48	3	Development workflow and methodology skills
Security	7	1	Security tooling skills (Prompt Security)
Vertical	86	4	Domain-specific: legal, biotech, DevOps, embedded

The company category includes skills from Microsoft (143), OpenAI (32), Sentry (16), HashiCorp (13), WordPress (13), Cloudflare (9), Expo (9), Hugging Face (9), Vue.js (8), Google (7), Better Auth (6), Vercel (5), Callstack (4), Tinybird (3), Neon (3), Black Forest Labs (2), Stripe (2), and others.

Each source repository is tracked as a git submodule pinned to a specific commit, ensuring reproducibility. Snapshot metadata (commit SHA, date, remote URL) is recorded for every source.

Skills were validated using `skill-validator` v0.4.0 (Carey 2026), a Go CLI tool that checks skills against the agentskills.io specification ("Agent Skills Specification" 2025).

2.2 Structural Validation

Each skill was evaluated against the specification's requirements:

- **Structure**: Presence of `SKILL.md`, correct directory layout, no unexpected files
- **Frontmatter**: Valid YAML frontmatter with required fields (name, description, version)

- **Content**: Markdown body completeness and formatting
- **Token budget**: Total context window usage

The validator produces a pass/fail result per skill along with categorized errors and warnings. We run the validator in two passes: first with structural checks only (`--only structure`), then with all other checks (`--skip structure`) for content analysis, contamination detection, and link validation. Link results are split into two categories: **internal links** (references to files within the skill directory, e.g. `references/api_guide.md`) are treated as structural failures since they indicate missing files, while **external links** (HTTP/HTTPS URLs) are reported separately as link health metadata. External URL validation is environment-dependent (URLs may be temporarily unreachable due to DNS failures, rate limiting, or transient outages), so excluding them from the pass/fail determination ensures reproducible structural compliance results across runs.

2.3 Content Analysis

We developed two automated content metrics:

Information density measures the proportion of a skill's content that consists of actionable material (code blocks and imperative sentences) versus prose. For skills containing code blocks, it is computed as:

$$\text{density} = 0.5 \times \frac{\text{code block words}}{\text{total words}} + 0.5 \times \frac{\text{imperative sentences}}{\text{total sentences}}$$

For prose-only skills (no code blocks), density equals the imperative sentence ratio alone, avoiding penalizing skills that legitimately contain no code. Imperative sentences are identified by checking whether the first word (after stripping markdown formatting) is a recognized imperative verb from a curated list of 46 common instruction verbs (e.g., "use", "create", "configure", "ensure", "avoid").

Instruction specificity measures the strength of a skill's language, based on the ratio of strong directive markers to weak advisory markers:

$$\text{specificity} = \frac{\text{strong markers}}{\text{strong markers} + \text{weak markers}}$$

Strong markers (10 patterns) include "must", "always", "never", "shall", "required", "do not", "don't", "ensure", "critical", and "mandatory". Weak markers (10 patterns) include "may", "consider", "could", "might", "optional", "possibly", "suggested", "prefer", "try to", and "if possible". All matching is case-insensitive with word-boundary constraints.

2.4 LLM-as-Judge Quality Scoring

To complement the heuristic content metrics, we employed an LLM-as-judge approach to evaluate skill quality across dimensions that resist automated measurement. The scoring uses a two-pass design:

Pass 1 — SKILL.md scoring evaluates each skill's primary instruction file on six dimensions (each scored 1–5): *clarity* (unambiguous instructions), *actionability* (step-by-step executability), *token efficiency* (conciseness), *scope discipline* (focus on stated purpose), *directive precision* (use of strong imperatives vs. vague suggestions), and *novelty* (information beyond LLM training data).

Pass 2 — Reference file scoring evaluates each reference file on five dimensions (each scored 1–5): *clarity*, *instructional value* (concrete examples and patterns), *token efficiency*, *novelty*, and *skill relevance* (alignment with the parent skill's purpose). The judge receives the parent skill's name and description as context.

Both passes use Claude Sonnet (claude-sonnet-4-5-20250929) with content truncated to 32,000 characters per document and results cached for reproducibility. Coverage is complete: all 673 skills were scored, along with 1,877 reference files across 411 skills. Three skills containing embedded AI-directed prompts caused persistent prompt interference with the judge model (see Limitations); these were scored manually.

2.5 Cross-Contamination Detection

Cross-contamination occurs when a skill designed for one context leaks information that influences agent behavior in another context. Research has established that code LLMs exhibit "Programming Language Confusion," systematically generating code in unintended languages despite explicit instructions, with strong defaults toward Python and shifts between syntactically similar language pairs (Moumoula et al. 2026). LLMs also exhibit a "copy bias" where they replicate patterns from in-context examples rather than reasoning independently about the task (Ali et al. 2024), and in-context code examples have been shown to bias the style and characteristics of generated code (Li et al. 2025). These findings suggest that mixed-language code examples in skill files could induce cross-language interference.

We developed a detection heuristic to estimate cross-contamination risk based on three structural factors:

1. **Multi-interface tool detection**: Does the skill reference a tool known to have multiple language SDKs (e.g., MongoDB, AWS, Docker)?
2. **Language mismatch**: Do code block languages differ from the skill's primary language category? Mismatches are weighted by syntactic similarity: mixing application languages (e.g., Python and JavaScript) carries higher weight than mixing an application language with auxiliary languages (shell, config, markup), reflecting research showing that Programming Language Confusion occurs primarily between syntactically similar language pairs (Moumoula et al. 2026).
3. **Scope breadth**: How many distinct technology categories does the skill reference?

These factors combine into a risk score from 0 to 1, classified as low (< 0.2), medium (0.2–0.5), or high (≥ 0.5). This score measures the *structural* potential for cross-contamination based on the presence of mixed-language content. We examine this structural metric against measured behavioral impact in the Behavioral Evaluation subsection of Cross-Contamination Risk.

An important distinction underlies the scoring: the research literature supports two different mechanisms by which multi-language content can degrade code generation, with different risk profiles:

- **Language confusion** — the model generates code using patterns from the wrong language. Research shows this primarily affects syntactically similar language pairs (C#/Java, JavaScript/TypeScript) and is driven by syntactic overlap in training data (Moumoula et al. 2026). Skills mixing multiple application-language SDKs (e.g., Python and JavaScript examples for the same API) carry the highest risk.
- **Context dilution** — additional content of any type consumes context window budget and reduces the model's attention to the user's actual task (Tian and Zhang 2025; Hong et al. 2025). This affects all multi-language content equally, regardless of syntactic similarity, and is addressed separately in our token budget analysis.

Many skills in our dataset mix an application language with bash scripts, YAML configuration, or SQL queries, a common and often necessary pattern for infrastructure and DevOps skills. Our scoring weights these auxiliary-language mismatches lower than application-to-application mismatches, since the syntactic dissimilarity between (for example) bash and Python makes language confusion less likely than between Python and JavaScript.

2.6 Behavioral Evaluation

To explore what mechanisms drive actual degradation in agent output when skills are loaded, and whether structural risk metrics capture them, we conducted a controlled behavioral evaluation of 19 skills from our dataset (one skill, doc-coauthoring, was excluded from aggregates because it triggers behavioral override rather than contamination: the model follows the skill's collaborative workflow instead of generating code).

Skill selection. We sampled 20 skills spanning the contamination score range (0.00–0.93), all source categories, and multiple content types. For each skill, we designed five tasks spanning five interference categories: *direct target* (tasks the skill is designed to help with), *cross-language* (tasks in a different language than the skill's examples), *similar syntax* (tasks in a syntactically similar language), *grounded* (tasks with well-defined correct answers the model already handles well), and *adjacent domain* (tasks in a related but distinct domain).

Task design. Tasks and expected patterns were designed with knowledge of each skill's content. For each skill, we identified the primary contamination vectors by reading the SKILL.md and reference files, then designed tasks to exercise those specific vectors (for example, requesting Go code for a skill whose examples are primarily in Python). Expected patterns (what correct code looks like) were derived from target-language SDK documentation; anti-patterns (what contaminated code looks like) were derived from the skill's non-target-language content. The LLM judge scored independently on four generic dimensions without knowledge of expected or anti-patterns. This design maximizes sensitivity to the specific interference mechanisms each skill could introduce, but means the eval functions as a targeted hypothesis confirmation test rather than a measurement of general-purpose degradation. We assess this limitation experimentally below (see "Probing Evaluation Sensitivity" in the Partial Knowledge section).

Conditions. Each task was evaluated under three conditions, each run three times at temperature 0.3:

- **Baseline (A)**: The task prompt alone, with no skill content
- **With-skill (B)**: The task prompt preceded by the skill's SKILL.md and selectively loaded reference files
- **Realistic (D)**: The task prompt preceded by the skill content, a Claude Code system preamble, and simulated conversation history, approximating how skills are loaded in practice

For skills with large reference directories, we loaded 2–3 reference files per task (selected for relevance to each task's interference vector) rather than all references. This selective loading approach was motivated by a controlled experiment showing that loading all reference files for a large skill produces near-identical outputs across runs, effectively reducing n to 1 and masking real interference effects.

Scoring. Outputs were scored by an LLM judge (Claude Sonnet) on five dimensions (correctness, completeness, code quality, specificity, following instructions) plus deterministic pattern matching for task-specific contamination signals. The composite score (1–5 scale) was averaged across dimensions and runs. Statistical significance was assessed using Welch's t-test on per-run composite scores.

Metrics. The primary metric is the B-A delta: mean with-skill score minus mean baseline score, measuring skill-only interference. The D-A delta measures interference under realistic conditions. The mitigation ratio (D-A / B-A) captures how much realistic context attenuates the skill-only effect.

3 Findings

3.1 Structural Compliance

Of 673 skills evaluated, **525 (78.0%) passed** structural validation and **148 (22.0%) failed**. Structural validation includes internal link integrity (references to files within the skill directory) but excludes external URL checks, which are environment-dependent and reported separately (see Methodology).
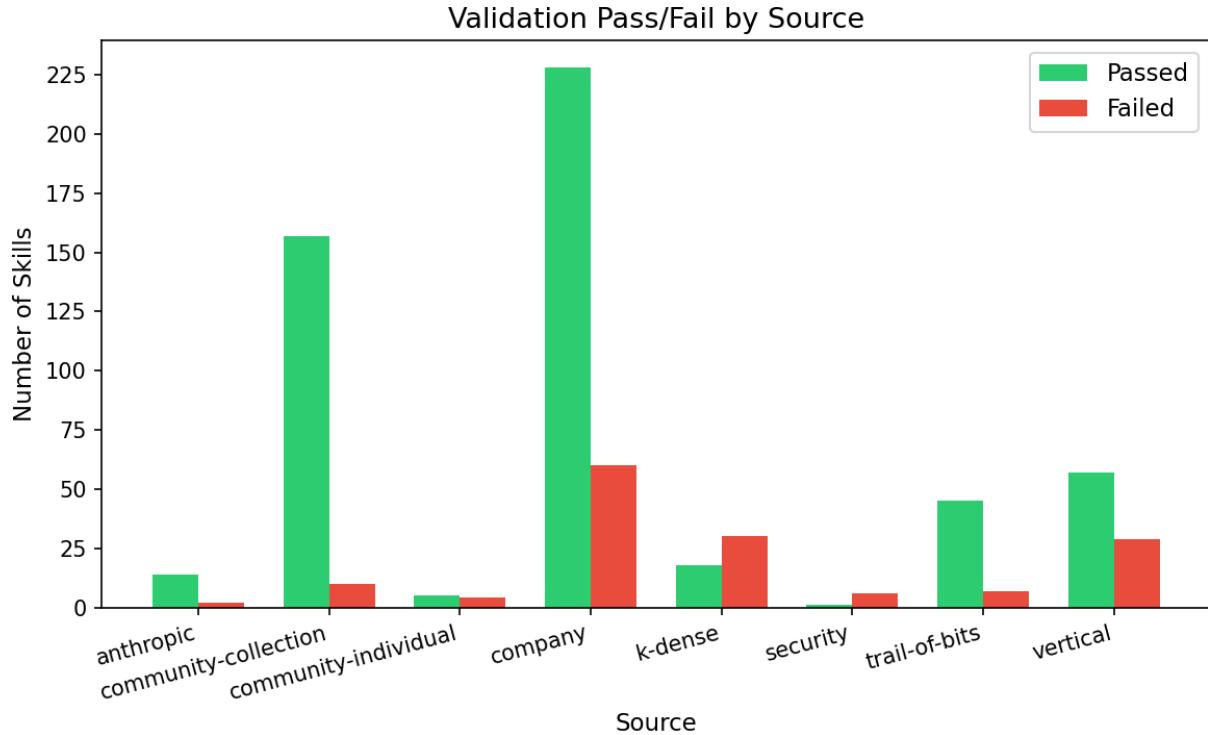
Figure 1: Pass/fail rates by source

Pass rates varied widely by source category:

Category	Skills	Pass Rate	Errors	Warnings
Community collections	167	94.0%	30	208
Anthropic	16	87.5%	7	37
Trail of Bits	52	86.5%	15	73
Company	288	79.2%	96	464
Vertical	86	66.3%	41	152
Community individual	9	55.6%	7	68
K-Dense	48	37.5%	62	116
Security	7	14.3%	6	23

Community collections lead at 94.0%, followed by Anthropic (87.5%) and Trail of Bits (86.5%). **Company-published skills have a lower pass rate (79.2%) than community collections (94.0%)**, inverting the common assumption that official company skills would be higher quality. The primary drivers:

- **Microsoft** (143 skills): Many skills use non-standard directory structures, placing skills under `.github/skills/` rather than the spec-standard layout. While functional within their GitHub Copilot integration, they generate structural validation errors.
- **Several companies** published skills before the spec was finalized and have not updated them to current requirements.

The K-Dense (Superpowers) skills had a low pass rate (37.5%), primarily because they use an alternative directory structure optimized for their development workflow rather than the agentskills.io specification.

Common errors across all sources included: - Missing or malformed YAML frontmatter - Unexpected files at the skill root (should be in `references/` or `assets/`) - Missing required frontmatter fields

3.1.1 Token Usage

Token counts varied by several orders of magnitude:
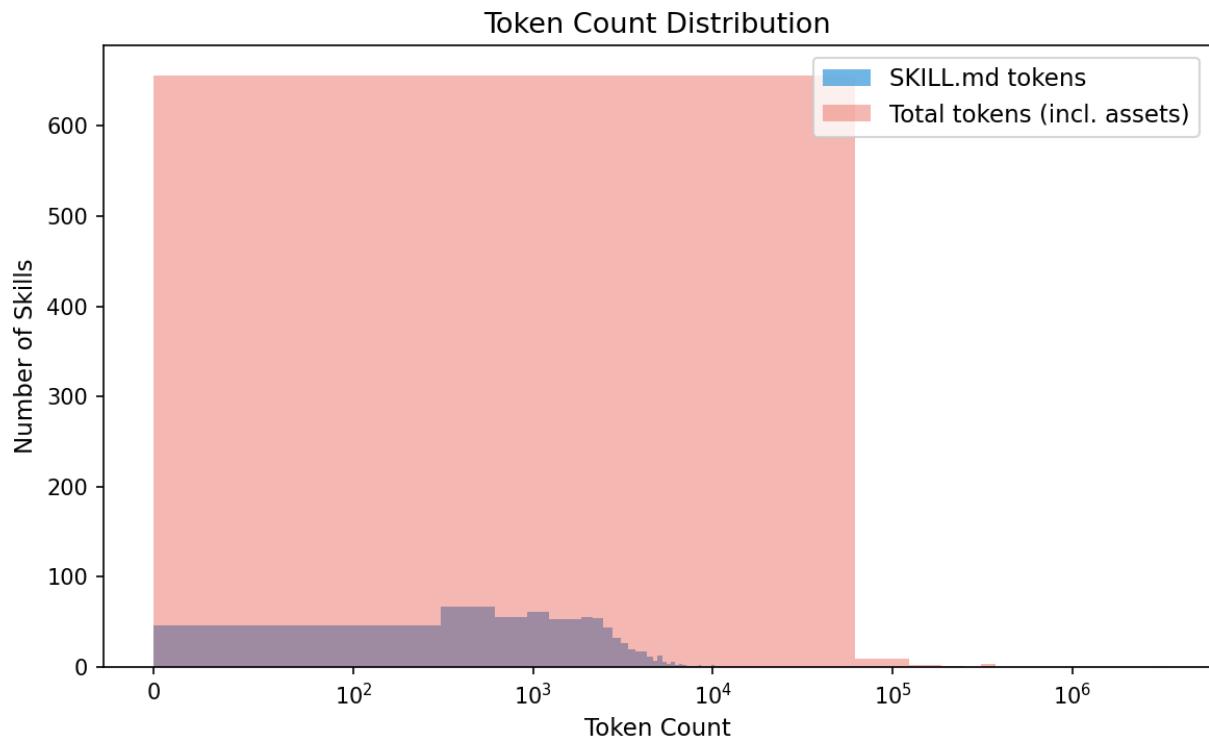
Figure 2: Token count distribution

- Minimum: 0 tokens (empty skill placeholders)
- Maximum: 3,098,484 tokens (a scientific computing skill with large reference datasets)
- Median: 5,236 tokens
- Mean: 16,711 tokens

Company-published skills tend to be more concise (average 6,790 tokens) than community collections (22,900 tokens). The most focused skills, from Anthropic (8,189 avg) and K-Dense (3,592 avg), demonstrate that effective skills can be compact.

3.1.2 Token Budget Composition

Breaking total token counts into their constituent parts (SKILL.md body, reference files, asset files, and nonstandard files) reveals a clear pattern: **52% of all tokens across the ecosystem are nonstandard files** that fall outside the specification's defined structure.
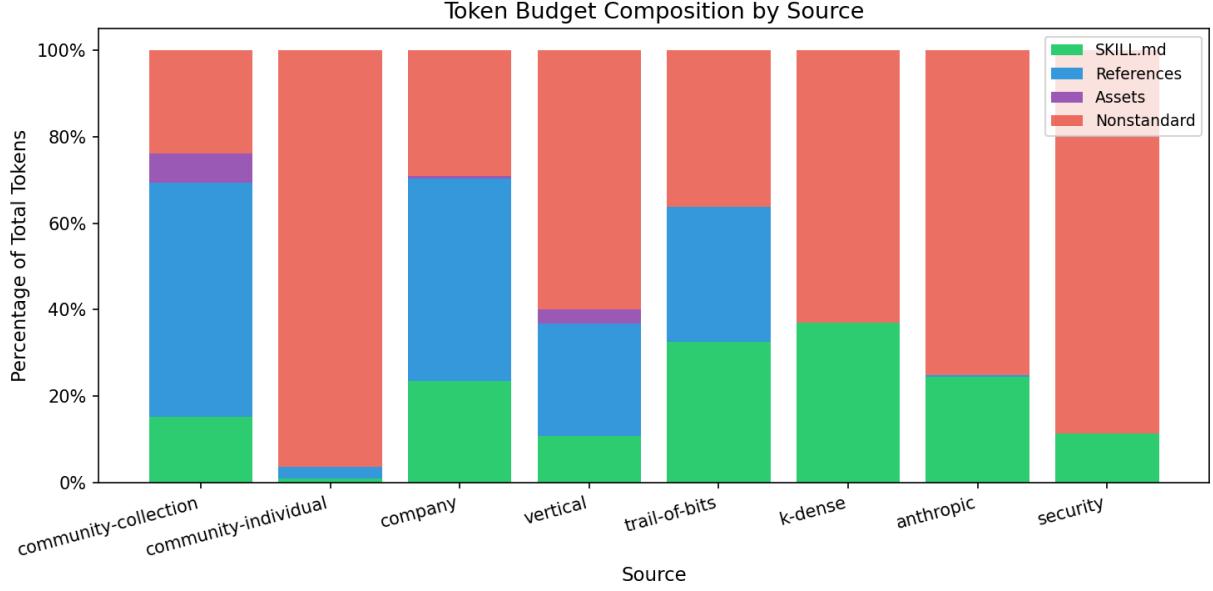

Figure 3: Token budget composition by source

The agentskills.io specification defines three categories of skill content: the primary `SKILL.md` instruction file, supplementary `references/` files, and binary `assets/` files. Any file outside these categories (placed at the skill root or in non-standard directories) may be loaded into the agent's context window depending on the platform's loading behavior. The specification does not explicitly prohibit loading nonstandard files, and agent platforms vary in whether they load the full skill directory or only recognized categories. When loaded, these files provide no instructional value, and research demonstrates that irrelevant context actively degrades LLM task performance (Shi et al. 2023; Liu et al. 2024). We term these **nonstandard tokens**.

Category	SKILL.md	References	Assets	Nonstandard
Community collections	15.2%	54.1%	6.8%	23.9%
Community individual	0.9%	2.6%	0.1%	96.4%
Company	23.5%	46.9%	0.6%	29.1%
Vertical	10.8%	26.0%	3.2%	60.0%
Trail of Bits	32.6%	31.1%	0.0%	36.3%
K-Dense	37.0%	0.0%	0.0%	63.0%
Anthropic	24.4%	0.4%	0.0%	75.1%
Security	11.4%	0.0%	0.0%	88.6%
Overall	**13.1%**	**32.1%**	**2.9%**	**52.0%**

Of 673 skills, 185 (27.5%) contain nonstandard files. The impact is substantial: the mean token count is inflated 108% by nonstandard files (mean effective tokens = 8,027 vs. mean total = 16,711). Even at the median, nonstandard files add 41% overhead.
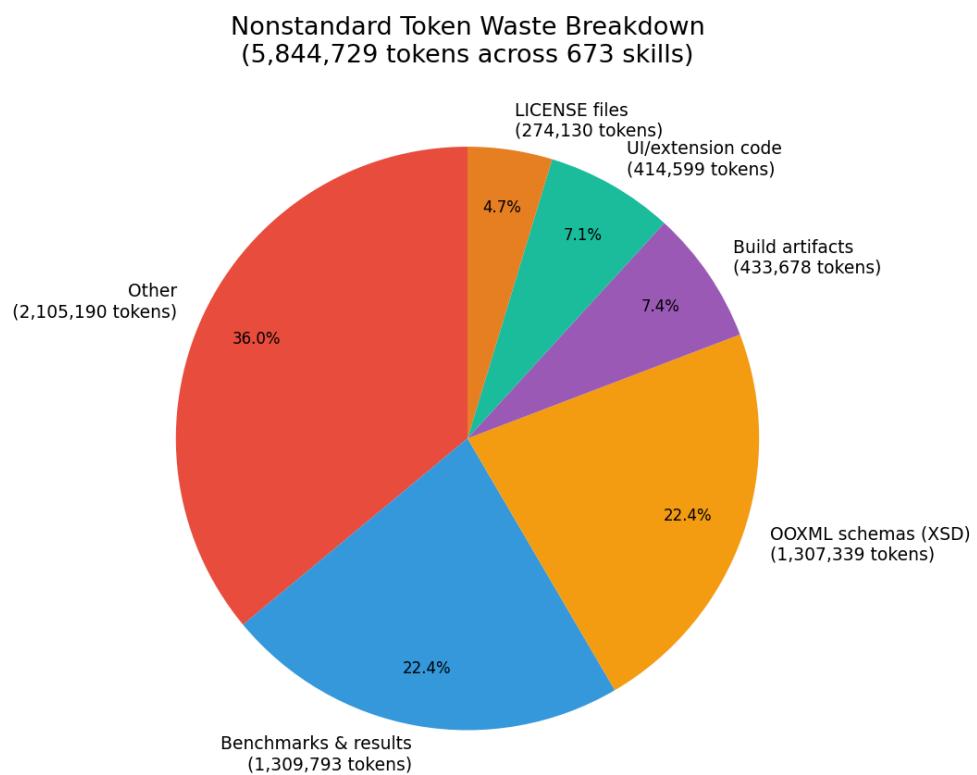
Figure 4: Breakdown of nonstandard token waste

Analyzing the 5.8 million nonstandard tokens reveals several distinct categories of waste:

- **OOXML schemas** (1.3M tokens, 22.4%): Four document-processing skills (docx, pptx, and their Anthropic variants) ship raw ISO-IEC 29500-4 XML Schema Definition files. A single schema file (`sml.xsd`, the SpreadsheetML schema) consumes 298,868 tokens, larger than most entire skills.
- **Benchmarks and results** (1.3M tokens, 22.4%): One skill (loki-mode) includes SWE-bench evaluation results and prediction logs, consuming over 500k tokens of JSON benchmark data.
- **Build artifacts** (434k tokens, 7.4%): Source maps (`.js.map`), lockfiles (`package-lock.json`), and compiled presentations (`.pptx`) that are development artifacts, not instructional content.
- **LICENSE files** (274k tokens, 4.7%): 89 skills include LICENSE.txt files at the skill root. While legally appropriate, each consumes ~2,700 tokens of context window space that provides no value to the agent.
- **UI/extension code** (415k tokens, 7.1%): VS Code extension source code and dashboard HTML/JavaScript shipped alongside skills.

The practical consequence is significant. An agent loading one of these skills receives a context window filled with XML schemas, benchmark JSON, or license text instead of the user's code. Research on LLM agent trajectories has shown that 40–60% of input tokens in agentic systems can be classified as useless, redundant, or expired information that is safely removable without performance loss (Xiao et al. 2025). For the 82 skills where SKILL.md represents less than 10% of total tokens, the instruction file risks being drowned out by supplementary material, consistent with findings that LLM performance degrades when relevant information is surrounded by irrelevant context (Liu et al. 2024; Shi et al. 2023).

Even Anthropic's own skills, the reference implementation from the spec authors, have the highest percentage of nonstandard tokens (75.1%) of any source category, driven primarily by LICENSE.txt files and template directories. This suggests the problem is structural rather than a matter of author diligence: the specification does not currently warn against or penalize nonstandard files, so authors have no signal that these files consume context window budget.

3.2 Content Quality

Content quality metrics revealed wide variation:

- **Information density**: Mean 0.206 (range 0.0–0.56). Most skills are prose-heavy with relatively few code examples or imperative instructions.
- **Instruction specificity**: Mean 0.616 (range 0.0–1.0). The expanded dataset shows a lower average specificity than our initial sample, driven by company skills that use more advisory language.

Anthropic skills cluster in the moderate-density, high-specificity quadrant: they are well-structured with clear directives but are not code-heavy. Company skills show the broadest distribution, ranging from highly specific API reference skills to vague best-practices guides.

3.2.1 LLM-as-Judge Quality Assessment

LLM-as-judge scoring reveals more about quality than heuristic metrics alone. Across all 673 skills, the global dimension means are: clarity 4.02, actionability 4.17, token efficiency 3.47, scope discipline 4.51, directive precision 3.58, and novelty 3.12. Skills are generally clear and well-scoped but often verbose and redundant with training data; token efficiency and novelty are consistently the weakest dimensions.

Source rankings by overall LLM score diverge from structural compliance rankings: company leads (3.97), followed by K-Dense (3.94), Trail of Bits (3.93), Anthropic (3.68), vertical (3.66), community collections
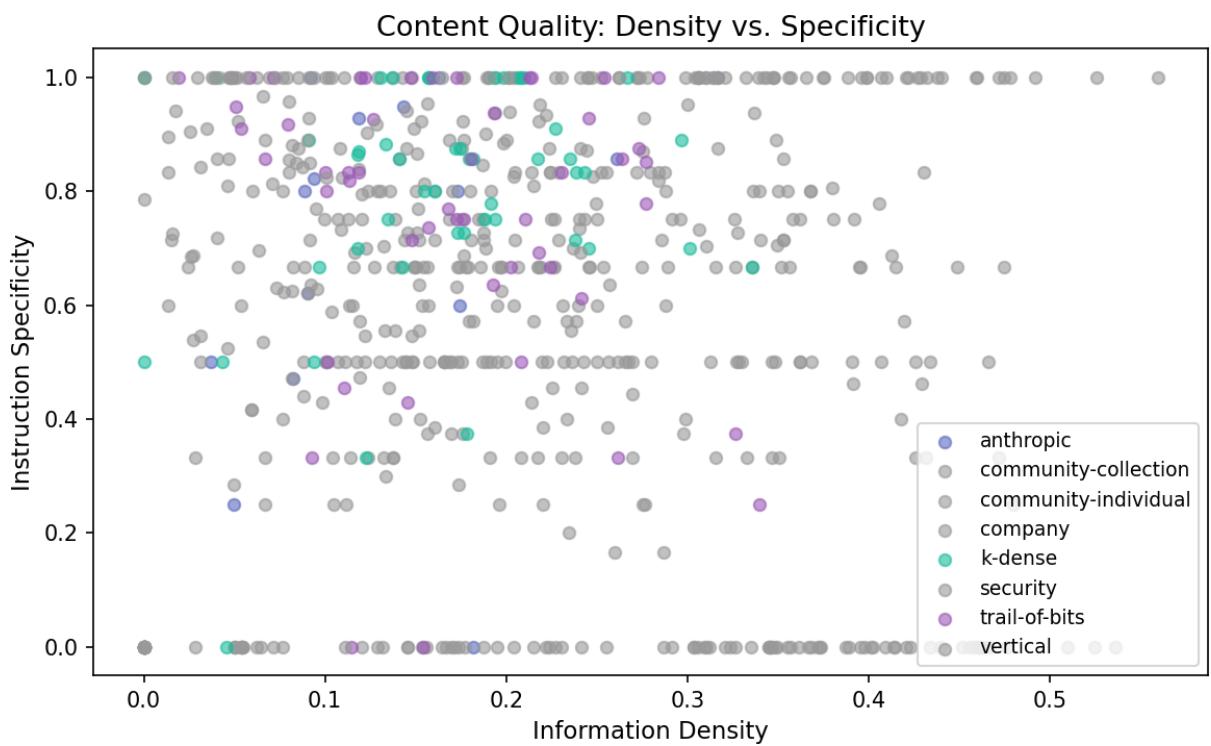
Figure 5: Information density vs. instruction specificity
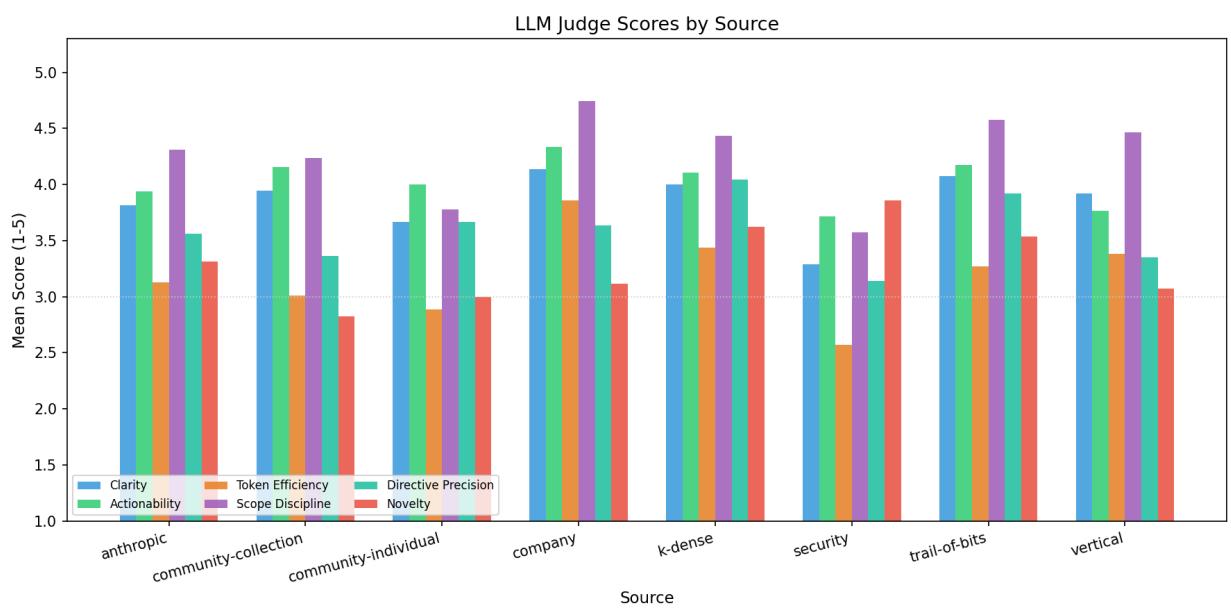
Figure 6: LLM judge scores by source category

(3.59), community individual (3.50), and security (3.36). Anthropic's ranking at #4 is notable: while their skills set the benchmark for instruction specificity (0.725), several experimental and template-like skills pull down the LLM average. The security category scores lowest, driven by skills with incomplete or scaffold-like content.

Structural validation is largely orthogonal to LLM-judged quality: skills that passed validation average 3.81 overall versus 3.80 for those that failed. A structurally valid skill is not necessarily a *good* skill, and vice versa; the two assessment methods measure different quality dimensions.

3.2.2   Quality by Source and Dimension

Disaggregating overall scores into individual dimensions reveals a craft-versus-content tradeoff across source categories. Token efficiency shows the largest spread across sources (1.29 points between company at 3.86 and security at 2.57), while novelty shows a comparable spread (1.03 points), indicating that sources differentiate on both writing quality and informational uniqueness.

Company-published skills exemplify the craft side: they rank #1 on scope discipline (4.74), clarity (4.13), actionability (4.33), and token efficiency (3.86), but drop to #5 of 8 sources on novelty (3.12). These skills are well-structured API documentation for tools that LLMs already know well. K-Dense skills are the mirror image: they lead on directive precision (4.04) and rank #2 on novelty (3.62), the two dimensions where company skills are weakest, but rank mid-pack on scope discipline. Their scientific computing methodology content is genuinely novel and uses strong imperative language.

Anthropic's relative strength is novelty (3.31, rank 4) despite ranking #6 on clarity (3.81), suggesting their skills prioritize unique content over polish. Community collections show the inverse: their relative strength is actionability (rank 3) but they are weakest on novelty (2.83, last among all sources), reflecting their tendency to provide step-by-step instructions for well-known libraries where the LLM already has strong coverage. Security skills are the floor on five of six dimensions but jump to rank 1 on novelty (3.86), since domain-specific security content is legitimately uncommon in training data, even when the skills themselves are poorly written.

The improvement path differs for each. Craft quality can be improved with better templates, linting, and editorial guidelines; it is a solvable problem. Novel domain knowledge, by contrast, is the irreducible value proposition of skills: information the LLM cannot derive from its training data. The sources that contribute the most unique information (K-Dense, Trail of Bits) are not the same sources that write the most polished skills (company publishers), suggesting that quality improvement efforts should target these two dimensions differently.

3.2.3   Novelty Analysis

Novelty, the degree to which a skill provides information beyond LLM training data, is the most discriminating dimension in our assessment. It measures whether skills offer genuine value that an agent cannot derive from its existing knowledge, making it arguably the most important quality signal for skill authors.

K-Dense and Trail of Bits lead on novelty, with 67% and 62% of their skills scoring 4 or above respectively. These sources focus on specialized domains (scientific computing methodologies, security vulnerability analysis) where proprietary workflows and non-obvious domain knowledge are central. Community collections trail at 28% scoring 4+, reflecting their tendency to wrap well-known libraries and frameworks where the LLM already has strong coverage.

Novelty correlates only weakly with the five craft dimensions ($r = 0.04$–0.39), confirming that it measures
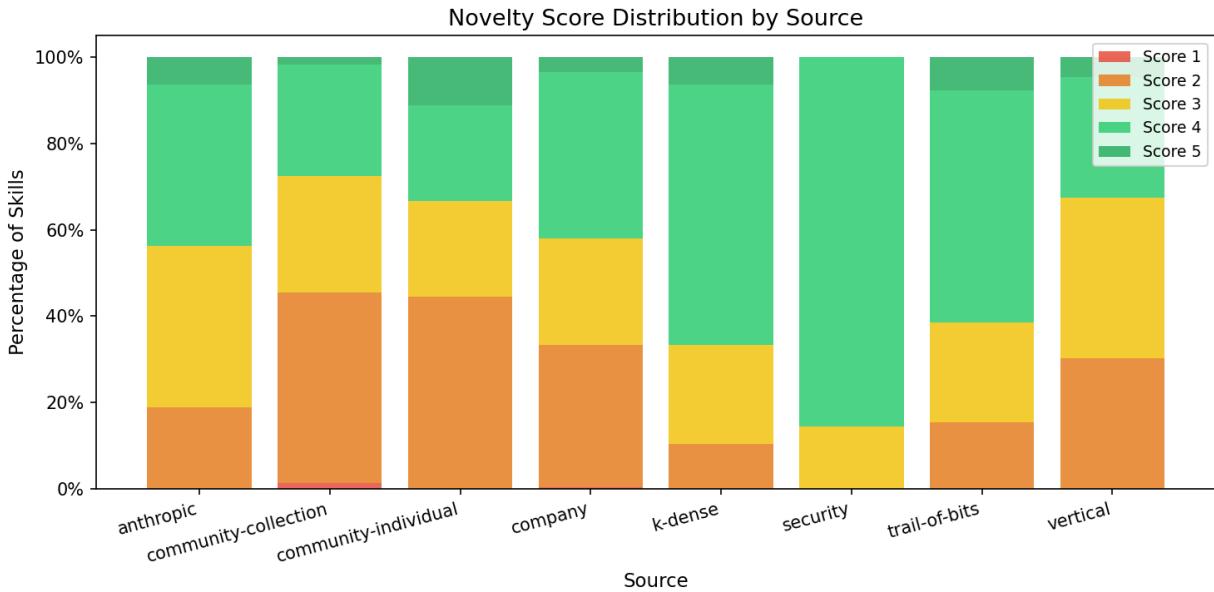
Figure 7: Novelty score distribution by source

something fundamentally different from writing quality. A skill can be clearly written, well-scoped, and concise (high craft scores) while still conveying information the LLM already knows (low novelty), a pattern common in community collection skills for popular frameworks.

3.2.4 Low-Novelty, High-Contamination Skills

Combining LLM-judged novelty with contamination risk scores reveals a structural pattern: **51 skills (7.6%) across the ecosystem have both low novelty (score ≤ 2) and medium-to-high contamination risk (score ≥ 0.2)**. These skills add mixed-language content with elevated structural complexity while providing little information the LLM does not already possess. In theory, this combination should produce a negative net effect on agent performance, though our exploratory behavioral evaluation (n = 19) found no correlation between structural contamination scores and behavioral degradation (r = 0.077), and the relationship between these structural indicators and actual impact remains an open question.

Company-published skills are disproportionately represented: 29 of 288 (10.1%) fall into this quadrant. Expanding the threshold to novelty ≤ 3 captures 45 company skills (15.6%). The most prominent cases are Azure SDK skills (`azure-identity-java`, `azure-security-keyvault-secrets-java`, and similar) with contamination scores above 0.50 but novelty scores of only 2. These are well-documented APIs whose documentation is heavily represented in LLM training data, packaged with multi-language examples that create the structural conditions for language confusion.

Novelty and contamination show near-zero correlation in our dataset (r = 0.00 for company skills, r = 0.07 across all skills, n = 673). There is no natural self-correction where contaminated skills compensate by being more novel. Among skills with medium or high contamination, company skills have the lowest mean novelty (3.12) compared to the non-company average (3.32); they are not making up for contamination risk with informational value.

The theoretical basis for concern is well-established: irrelevant context degrades LLM task performance (Shi et al. 2023), mixed-language content causes systematic programming language confusion (Moumoula

et al. 2026), and in-context examples bias generated code toward the patterns they contain (Ali et al. 2024). A skill that introduces these interference risks without contributing novel information would, by this reasoning, be worse than no skill at all: the agent pays the contamination cost without receiving informational benefit. However, our exploratory behavioral evaluation (n = 19) complicates this prediction: we observed a suggestive pattern where novelty correlates with the *magnitude* of behavioral effects in both directions (r = +0.267 for |B-A|, p ~ 0.27), and low-novelty skills showed smaller degradation than expected, as if the model largely ignores content it already knows. If this pattern holds in larger samples, the real risk may be high-novelty skills loaded for mismatched tasks, where the model pays attention to genuinely novel content that happens to be irrelevant to the current task. Confirming whether low-novelty/high-contamination skills actually degrade performance in practice requires behavioral testing of this specific subpopulation.

3.3 Cross-Contamination Risk
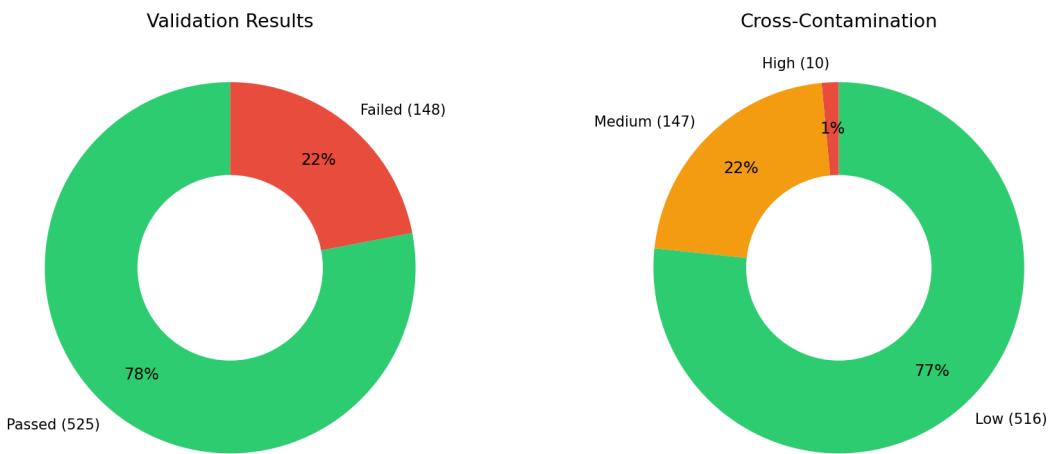
Figure 8: Validation and contamination overview

Our cross-contamination analysis identified:

- **10 high-risk skills** (1.5%) — substantial structural potential for cross-language interference, driven primarily by application-to-application language mixing
- **147 medium-risk skills** (21.8%) — some multi-language or multi-technology mixing
- **516 low-risk skills** (76.7%) — focused on a single technology or language

These risk levels measure language complexity (the structural presence of mixed-language content) rather than demonstrated behavioral impact. Our exploratory behavioral evaluation (see Behavioral Evaluation Results) found no correlation between these structural scores and measured degradation (r = 0.077, n = 19), suggesting that content-specific factors may matter more than language mixing per se.

The security category had the highest average risk score (0.344), followed by Trail of Bits (0.151) and company-published skills (0.127). Security tools inherently operate across multiple languages and environments, making this expected. Company skills, particularly those for cloud platforms (Azure, AWS, Terraform), scored higher because they often cover multiple language SDKs within a single skill.
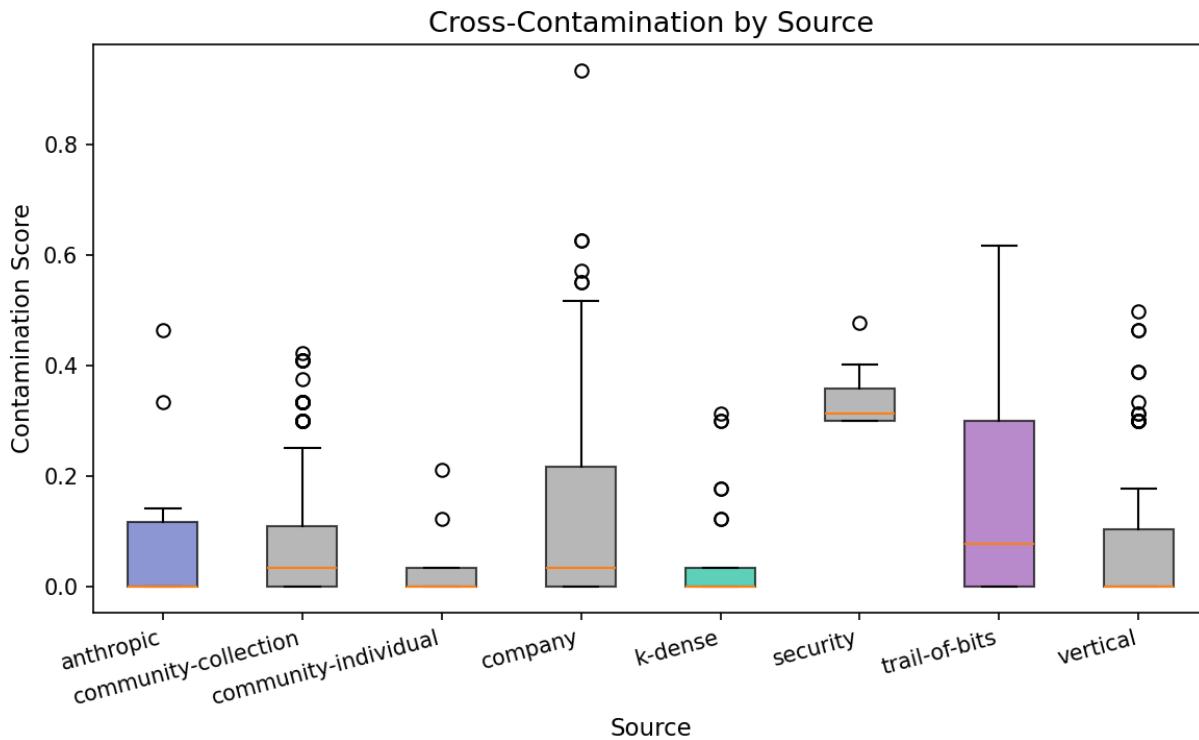
Figure 9: Contamination scores by source

3.3.1 High-Risk Skills

The 10 high-risk skills cluster around two primary patterns, with the similarity-weighted scoring concentrating high-risk flags on skills with genuine application-to-application language mixing:

1. **Multi-SDK platform skills** (highest language confusion risk): Azure, AWS, and Terraform skills that cover multiple application language bindings (Python, Java, TypeScript, .NET, Rust) within one skill. These represent the strongest contamination concern because they mix syntactically similar languages for the same API, exactly the pattern that Programming Language Confusion research identifies as most problematic (Moumoula et al. 2026).

2. **Infrastructure-as-code skills** (moderate risk): Skills mixing shell commands, configuration languages (HCL, YAML), and application code. The language confusion risk here is lower due to syntactic dissimilarity between auxiliary and application languages, though context dilution remains a concern.

3. **Security analysis skills** (mixed risk): Tools that inherently operate across language boundaries, combining both application-to-application mixing (higher risk) and application-to-auxiliary mixing (lower risk).

Key examples: - **upgrade-stripe** (Stripe, risk: 0.93): Covers SDK upgrades across multiple application languages (Python, Ruby, JavaScript) with mixed code examples. This is the highest-risk pattern, as examples for the same Stripe API in multiple syntactically similar languages create direct confusion potential - **copilot-sdk** (Microsoft, risk: 0.63): Multi-SDK skill mixing application languages with shared API patterns - **provider-resources** (HashiCorp, risk: 0.55): Terraform provider development mixing Go with HCL and shell. The Go application code mixed with infrastructure references drives the score, while the HCL/shell

17

auxiliary mixing is appropriately down-weighted - **ossfuzz** (Trail of Bits, risk: 0.53): Combines Docker, shell, Python, and C/C++ for fuzz testing. The Python/C++ application language mixing is the primary risk factor

monitoring-observability (DevOps vertical), previously scored as high-risk (0.72) under uniform weighting, now scores as medium (0.50) because its multi-language content is primarily bash, YAML, and configuration alongside infrastructure tool references, not application-to-application language mixing. This reclassification better reflects the actual language confusion risk.

3.3.2 Case Study: MongoDB Cross-Contamination

During development of this analysis, we observed an illustrative case of cross-contamination: an unpublished MongoDB skill containing `mongosh` (shell) examples caused Claude Code to generate incorrect Node.js driver code. The agent produced queries using shell syntax instead of the Node.js driver API, and it embedded shell-specific operators in JavaScript contexts.

To validate this observation experimentally, we constructed a controlled A/B eval using a MongoDB Search skill under development by the first author.[1] Five tasks were generated under baseline (no skill) and with-skill conditions (3 runs each at temperature 0.3, scored by an LLM judge and deterministic pattern matching). Two tasks produced clear contamination signals:

- **Shell syntax in JSON Schema output**: A task requesting a pure JSON Schema document produced `db.` (mongosh shell prefix) references in 0/3 baseline runs but 3/3 with-skill runs; the skill's reference files contain `mongosh` examples that bleed into non-shell output contexts.
- **Invalid JSON constructs in index definitions**: A task requesting valid JSON index definitions produced `ISODate()` function calls (a mongosh-specific construct invalid in JSON) in with-skill runs but not in baseline runs, alongside `//` comments (also invalid in JSON) sourced from the skill's reference examples.

Overall, the skill slightly degraded output quality: baseline averaged 4.37/5.0 across judge dimensions versus 4.25/5.0 with the skill loaded. The model already knows MongoDB Search well from training data; the skill's primary effect was introducing shell-syntax contamination rather than filling knowledge gaps. The realistic context condition (skill loaded alongside a Claude Code system preamble and simulated conversation history) mitigated the contamination in some tasks, a pattern that generalizes across the broader behavioral evaluation (see Behavioral Evaluation Results).

This is consistent with documented LLM behaviors: code LLMs exhibit Programming Language Confusion, systematically defaulting to patterns from syntactically similar languages (Moumoula et al. 2026), and in-context examples bias the style of generated code toward reproducing the patterns present in the examples (Li et al. 2025). The MongoDB case is particularly illustrative because the shell examples are syntactically valid JavaScript (MongoDB's Shell is JavaScript-based), making the interference subtle: the generated code *looks* correct but uses the wrong API for the target context. Research on attention dilution in code generation further suggests that as skill file content grows, the model pays less attention to the user's actual intent (Tian and Zhang 2025).

---

[1]Unlike the published skills analyzed elsewhere in this paper, the MongoDB Atlas Search skill was unpublished at the time of testing. The skill version tested, eval pipeline, task definitions, generation outputs, and scoring results are available at https://github.com/dacharyc/mdb-skill-builder/tree/main/eval to enable independent replication. The skill has since been revised based on the contamination findings described here.

3.3.3 Case Study: Gemini API Multi-Language Skill

The **gemini-api-dev** skill (Google, risk: 0.55) provides a concrete illustration of the API shape differences that make cross-contamination hard to detect. The skill demonstrates a single operation, `generateContent`, in four languages, each with a subtly different API shape:

Python — keyword arguments, snake_case method:

```python
response = client.models.generate_content(
    model="gemini-3-flash-preview",
    contents="Explain quantum computing"
)
print(response.text)
```

JavaScript — options object, camelCase method:

```javascript
const response = await ai.models.generateContent({
  model: "gemini-3-flash-preview",
  contents: "Explain quantum computing"
});
console.log(response.text);
```

Go — positional parameters, PascalCase method, explicit context:

```go
resp, err := client.Models.GenerateContent(ctx,
    "gemini-3-flash-preview",
    genai.Text("Explain quantum computing"), nil)
```

Java — positional parameters, camelCase method, null config:

```java
GenerateContentResponse response =
    client.models.generateContent(
        "gemini-3-flash-preview",
        "Explain quantum computing",
        null);
System.out.println(response.text());
```

The differences are subtle but consequential: Python uses keyword arguments while Java uses positional parameters; Go wraps content in `genai.Text()` while others pass raw strings; Java accesses the result via a method call (`response.text()`) while Python and JavaScript use a property (`response.text`). An agent working on a Python project but primed with the JavaScript or Go patterns from this skill might generate code using positional arguments, or omit the keyword parameter names, producing code that would fail at runtime. This is exactly the kind of syntactically plausible but semantically incorrect output that Programming Language Confusion research predicts for syntactically similar languages sharing the same API surface (Moumoula et al. 2026).

3.3.4 Behavioral Evaluation Results

To explore what mechanisms drive behavioral degradation when skills are loaded, we evaluated 19 skills using the methodology described in the Behavioral Evaluation subsection of Methodology. In this sample, **we found no correlation between structural contamination scores and behavioral degradation** (r = 0.077, n = 19). Skills with high structural risk scores did not consistently produce worse output than skills with low scores. Given the small sample size and targeted evaluation design, this is an exploratory finding,

but the individual case studies below illustrate why content-specific mechanisms may matter more than structural language mixing.
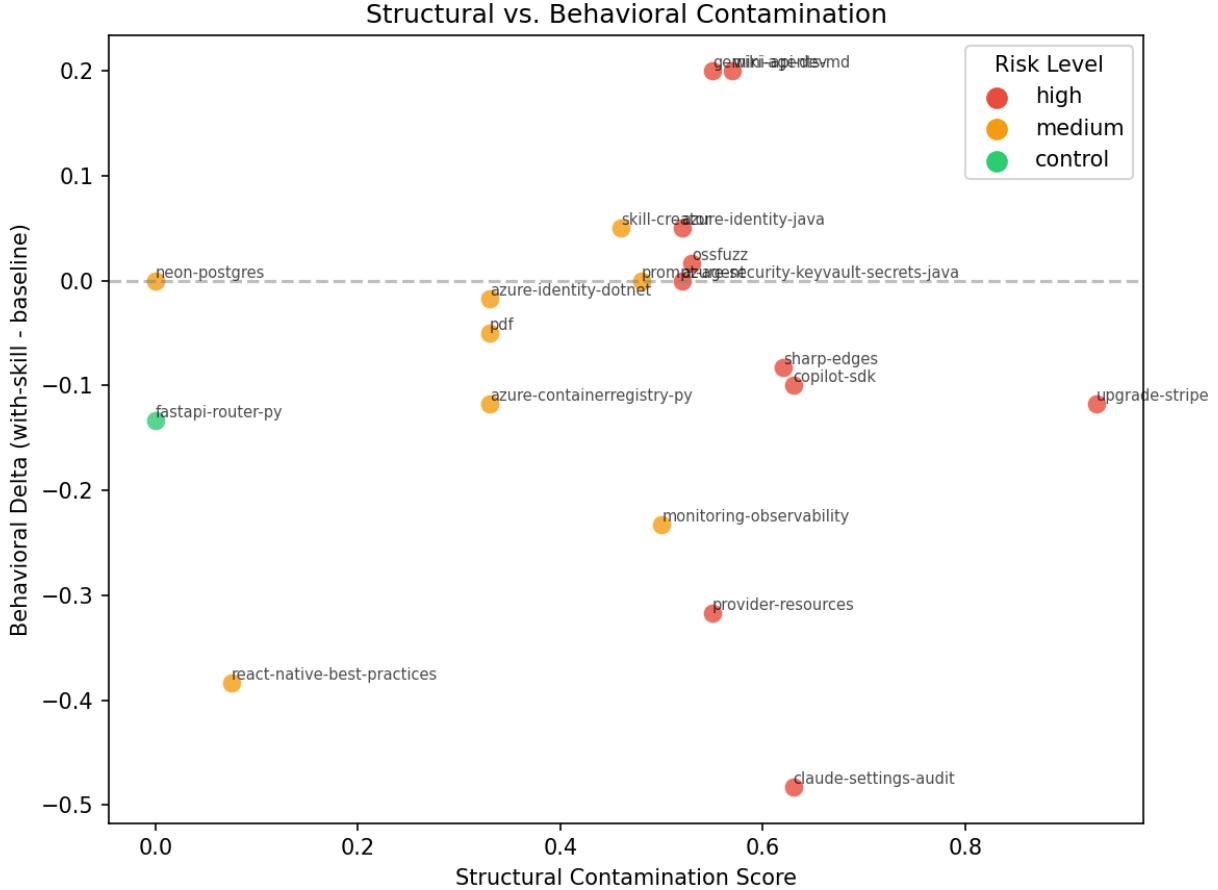
Figure 10: Structural contamination score vs. behavioral delta

The behavioral eval does find real degradation (mean B-A delta is -0.080 across 19 skills, with 6 skills reaching statistical significance at $p < 0.05$), but the degradation is driven by content-specific mechanisms rather than the structural language-mixing patterns our contamination heuristic measures. The disconnect is sharp in individual cases: **upgrade-stripe** (highest structural risk at 0.93) shows only modest behavioral degradation (B-A = -0.117, not significant), while **react-native-best-practices** (structural risk 0.07, the second-lowest in our sample) produces the largest degradation (B-A = -0.383, p = 0.001).

3.3.5 Content Interference Mechanisms

Examining the degradation patterns across all 19 skills reveals that cross-language code confusion, the mechanism our structural scoring was designed to detect, is only one of several interference vectors. We identified six distinct mechanisms:

1. **Template propagation**: A skill's output templates are reproduced verbatim in unrelated contexts. The **claude-settings-audit** skill (B-A = -0.483, p = 0.019) contains `// comments` in JSON output templates; the model faithfully reproduces this invalid JSON syntax across all tasks. The contamination is format-level, not language-level.
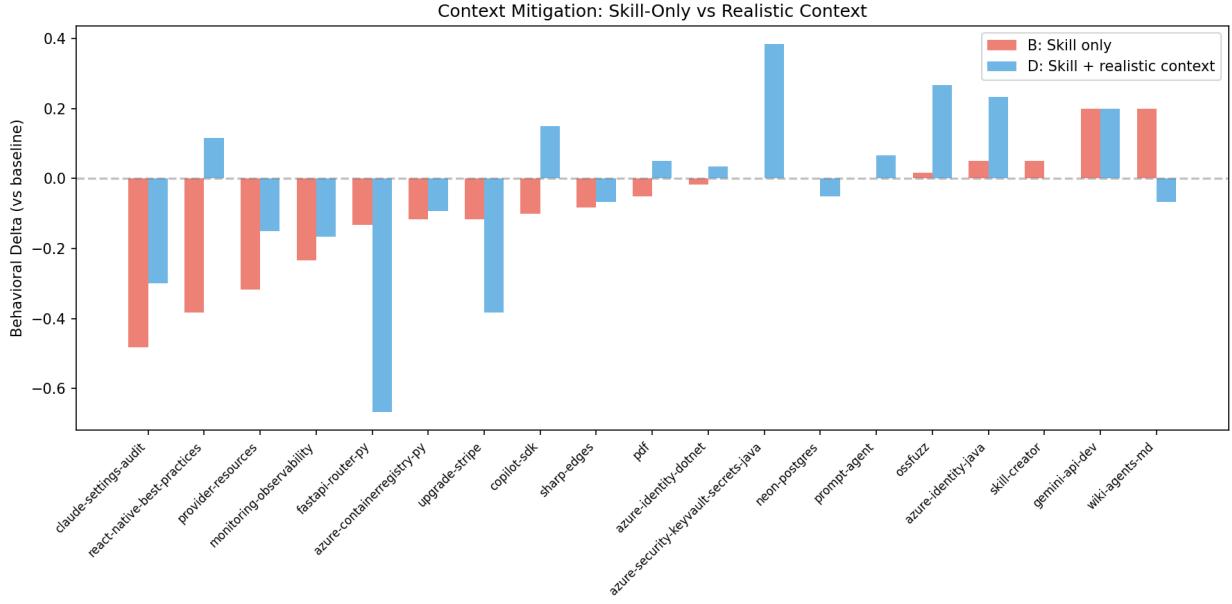
Figure 11: Behavioral context mitigation

2. **Textual frame leakage**: A skill's identity bleeds into output prose without affecting code correctness. **react-native-best-practices** outputs include phrases like "following React Native best practices for native iOS code" in tasks requesting pure Swift or Kotlin; the skill's framing contaminates the explanation even when the code is clean. This is scored as degradation by both the LLM judge and human review, but is a different phenomenon than language confusion.

3. **Token budget competition**: Skill content causes the model to generate longer outputs that hit token limits, truncating implementations. This takes two forms: in the react-native case, with-skill outputs spend 2.6× more tokens on explanatory prose and 18% fewer on code compared to baseline. In the **provider-resources** case (B-A = -0.317, p = 0.010), the skill's detailed Go patterns cause the model to generate 2× longer implementations (adding full import blocks, extra struct fields, and elaborate helper functions) that hit the 4,096-token ceiling mid-function. Baseline outputs for the same task use ~1,800 tokens and complete successfully; with-skill outputs expand to ~3,900 tokens and are truncated.

4. **API hallucination**: The model invents plausible but nonexistent API methods after seeing similar APIs in skill content. **upgrade-stripe** (B-A = -0.117) generates fabricated Stripe SDK methods that follow the naming conventions in the skill's examples but do not exist. Unlike language confusion, the hallucinated code is in the *correct* language; it is the API surface that is wrong.

5. **Cross-language code bleed**: The classic Programming Language Confusion mechanism: shell syntax appearing in JavaScript output (MongoDB case study), mongosh operators in JSON contexts. This is the only mechanism our structural scoring is designed to detect, and it does predict it: the MongoDB skill's structural score correctly flags the risk. But across 19 skills, this mechanism accounts for a minority of the total degradation observed.

6. **Architectural pattern bleed**: A skill's design-level patterns transfer across languages without any syntax errors. The provider-resources skill teaches Go Terraform provider conventions (separate models, client abstractions, test directories); when the model is asked to write a Python Pulumi provider, it generates an over-engineered 7-file project structure instead of a single-file implementation, consum-

21

ing tokens on scaffolding before reaching the core task. The judge flags this as "enterprise Java/C# patterns" because the Go architecture doesn't manifest as Go *syntax* in Python (which would be classic language confusion) but as inappropriate *structural* complexity. This is the subtlest mechanism we observe: the output is syntactically correct in the target language, but the design is contaminated.

These mechanisms have different implications for skill authors. Template propagation and API hallucination are addressable through content review (fix invalid syntax in templates, avoid suggestive API patterns). Textual frame leakage, token budget competition, and architectural pattern bleed are harder to mitigate because they emerge from the skill's identity, scope, and structural conventions rather than specific content defects.

3.3.6   Case Study: react-native-best-practices

The **react-native-best-practices** skill is the most instructive case because it produces the largest behavioral degradation (B-A = -0.383) despite having essentially zero structural contamination risk (0.07). The skill's 281 code blocks are 94% language-labeled, spanning 8 application languages (C++, JavaScript, JSX, Kotlin, Objective-C, Swift, TSX, TypeScript). By every structural metric, it should be well-organized.

The degradation comes from two interacting mechanisms. First, textual frame leakage: the skill's React Native framing appears in output prose for tasks that have nothing to do with React Native. A task requesting a native iOS media player in Swift receives an introduction referencing "React Native best practices for native iOS development." A task requesting a Kotlin Android service references "React Native's threading model." The code itself is largely correct; it is the surrounding explanation that is contaminated.

Second, token budget competition: with the skill loaded, outputs allocate substantially more tokens to explanatory commentary and less to code. Under the eval's 4,096-token output ceiling, this means with-skill outputs produce less complete implementations than baseline outputs for the same tasks.

A revealing contrast is **sharp-edges** (Trail of Bits): 12 distinct application languages, 282 code blocks, 100% language-labeled, structural contamination score 0.62, yet only -0.083 B-A delta. The difference is content type, not language count. sharp-edges teaches security vulnerability patterns that are conceptually portable across languages (buffer overflows look similar in C, C++, and Rust). react-native-best-practices teaches framework-specific implementation patterns (FlashList, Turbo Modules, threading models) tightly bound to specific platforms. The security patterns don't interfere; the framework patterns do.

3.3.7   Realistic Context Mitigates Most Interference

In our 19-skill sample, realistic context (a system preamble plus simulated conversation history, approximating how skills are actually loaded in agentic workflows) substantially attenuated interference. Mean D-A delta (realistic minus baseline) was -0.023, compared to mean B-A delta of -0.080. The mean mitigation ratio was -74.5%, suggesting that realistic context may eliminate a substantial portion of skill-only degradation. However, this ratio is computed from small, often non-significant deltas (for the 13 skills where B-A did not reach significance, the ratio is effectively noise divided by noise), so the precise magnitude should be interpreted cautiously.

This mitigation is consistent across most skills but not universal, and its effectiveness depends on the mechanism. **upgrade-stripe** is the clearest exception: its D-A delta (-0.383) is *worse* than its B-A delta (-0.117), meaning realistic context amplifies rather than mitigates the interference. The mechanism is API hallucination: in realistic context, the system preamble's emphasis on using available tools and being helpful appears to reinforce the model's tendency to generate plausible-but-fabricated Stripe API methods. We validated this finding by investigating a potential codebase snippet confound: the realistic condition injects a simulated codebase file read, and the default Python snippet used async patterns (FastAPI + AsyncSession) that
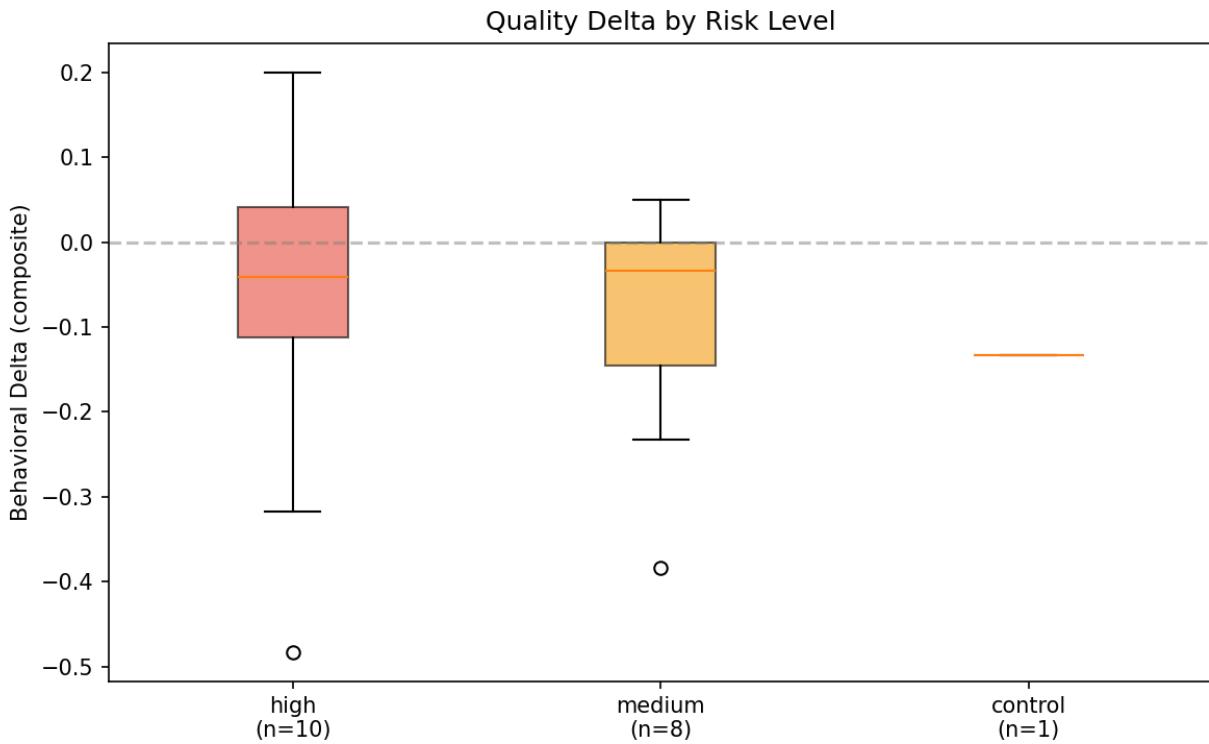
Figure 12: Behavioral deltas by risk level

could cause the judge to flag async/sync confusion as contamination. Re-running upgrade-stripe's Python task with a synchronous codebase snippet reduced that task's D-A delta from -0.833 to -0.250 (a ~70% reduction), confirming the confound. However, the skill-level amplification persists (D-A = -0.383 vs. B-A = -0.117) across the remaining tasks, and the judge flags a different set of fabrication signals (invented exception classes, misapplied method attributes) that are unambiguously skill-induced. The codebase snippet confound is relevant to other skills using language-specific snippets in Condition D, but the upgrade-stripe amplification pattern is robust to this correction. **provider-resources** shows an asymmetry within the same skill: cross-language interference (Go patterns in Python output) is fully mitigated by realistic context, but same-language over-specification (Go patterns in a different Go task) shows only 22% mitigation. When the skill's patterns are in the same language as the task, the model has less reason to discard them.

If this pattern holds at scale, it would be reassuring for skill authors and platform maintainers: the degradation measured under the artificial skill-only condition (B) may substantially overestimate the interference that users experience in practice. Skills are not loaded in isolation; they compete with a rich context that anchors the model's behavior.

3.3.8 Exploring Predictors of Behavioral Degradation

If structural contamination scores show no correlation with degradation in our sample, what does? The following correlations are exploratory: with n = 19 skills, none would survive multiple-comparison correction, and all should be treated as hypotheses for future testing rather than established findings.

Novelty amplification (exploratory). The largest correlation we observed with degradation *magnitude* is the skill's novelty score (r = +0.267 for |B-A|, n = 19, p ~ 0.27). If this pattern holds in larger samples, it

would suggest that high-novelty skills produce larger behavioral effects in both directions, helping more on tasks they're designed for and hurting more on tasks they're not. This would be consistent with the model paying more attention to genuinely novel content. We also observed suggestive task-type interactions: on cross-language tasks, higher novelty showed a negative association with performance (r = -0.218); on similar-syntax and adjacent-domain tasks, novelty showed a positive association (r = +0.417, +0.344). These per-task-type correlations are computed on even smaller subsets and should be interpreted with particular caution.

Task type. Degradation was concentrated on grounded tasks (mean B-A = -0.193) and cross-language tasks (mean B-A = -0.171), while direct-target and adjacent-domain tasks showed near-zero mean effect (-0.009 each). This pattern, where skills tend to help on their intended domain and hurt on domains where the model already performs well without them, is the most robust descriptive finding from the behavioral eval, as it does not depend on cross-skill correlations.
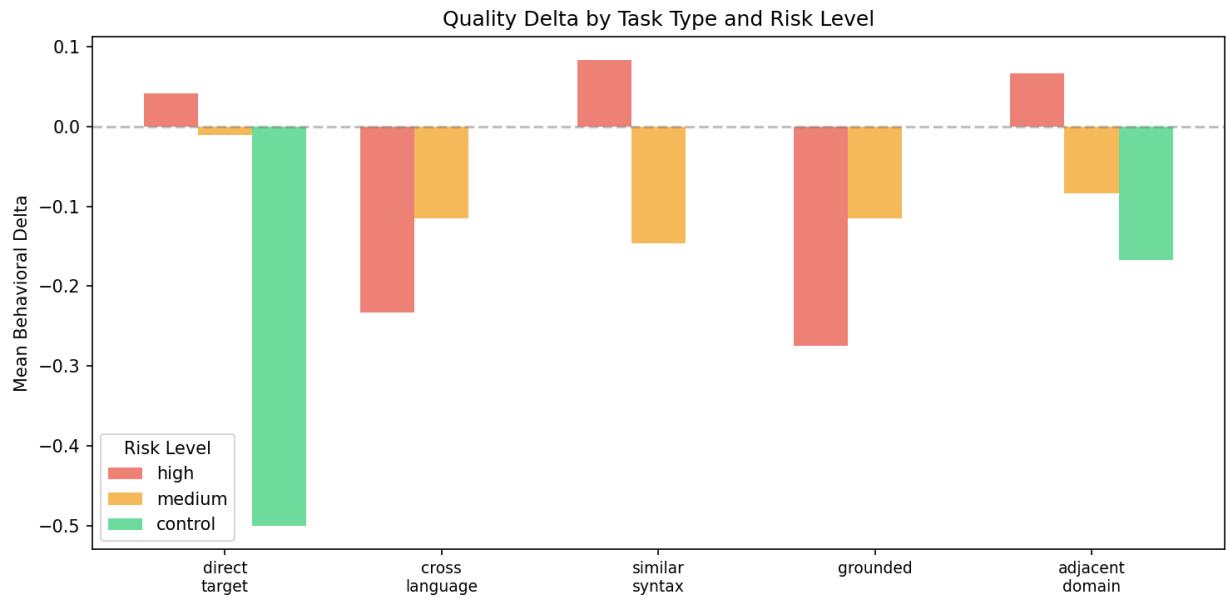

Figure 13: Behavioral task type effects

Negative results. Several factors we expected to correlate with degradation did not in our sample:

- *Code block language labels*: Skills with 100% language-labeled code blocks showed *worse* mean degradation (-0.105) than partially-labeled skills (-0.052), the opposite of what PLC research would predict. Labels address within-context language disambiguation, not the cross-task interference mechanisms we observe.
- *Skill size*: Token overhead correlated weakly with B-A (r = -0.188), but the relationship is confounded by measurement artifacts from selective reference loading.
- *LLM quality dimensions*: Most individual quality dimensions (clarity, actionability, token efficiency, scope discipline) showed weak correlation with B-A (all |r| < 0.15). Directive precision was the exception (r = +0.36), suggesting that skills with more imperative language may produce slightly larger signed deltas, though this does not survive multiple-comparison correction. Directive precision also shows the strongest correlation with novelty among the craft dimensions (r = 0.39 vs. r = 0.04–0.20 for the others), so this effect is likely confounded: skills that teach genuinely novel conventions tend to use more imperative language ("always use X", "never do Y"), and the behavioral correlation may reflect the underlying novelty rather than the directive style itself.

24

- *Structural contamination score*: r = 0.061, essentially zero.

The practical implication is that contamination risk is unlikely to reduce to a single structural metric. The case studies above suggest that content specificity (framework-specific patterns vs. language-portable patterns), task mismatch (skill loaded for an unrelated task), and content defects (invalid syntax in templates, suggestive API patterns) may matter more than language mixing per se, but confirming this requires a larger behavioral sample.

3.3.9   Partial Knowledge and API Fabrication

The API hallucination mechanism (mechanism 4 above) suggests a specific structural risk factor related to incomplete API documentation. Three skills in our evaluation produce the most fabrication-type judge signals, where the judge flags nonexistent methods, invented exception classes, or fabricated version numbers:

Skill	With-skill signals	Fabrication signals	Reference files	SKILL.md content
copilot-sdk	84	24 (29%)	None	15K chars, pattern-level API guide
upgrade-stripe	12	5 (42%)	None	5.6K chars, migration/versioning patterns
claude-settings-audit	36	3 (8%)	None	11K chars, audit workflow templates

All three provide pattern-level API knowledge (method naming conventions, architectural patterns, versioning strategies) without comprehensive SDK documentation. By contrast, skills with substantial reference files (azure-identity-java: 37K total, neon-postgres: 88K, sharp-edges: 127K, react-native-best-practices: 161K) show near-zero fabrication signals in the with-skill condition, even when they exhibit other contamination types (cross-framework patterns, textual frame leakage).

The copilot-sdk case is partially confounded: the model hallucinates Copilot SDK methods even at baseline (23 fabrication signals without any skill), reflecting post-training knowledge gaps rather than skill-induced fabrication. But upgrade-stripe shows the clearest skill-induced pattern: 0 fabrication signals at baseline, 5 with the skill loaded, 9 under realistic context. The skill's migration vocabulary (version pinning, per-request overrides, SDK-specific exception hierarchies) gives the model enough conceptual scaffolding to construct plausible-sounding but incorrect API calls (fabricated version strings like `stripe-go/v81` and `2026-01-28.clover`, nonexistent Go constants like `stripe.ErrorTypeConnection`, and incorrect method signatures like passing `toleranceSeconds` as a positional argument to `constructEvent`).

This pattern is consistent with a **partial knowledge hypothesis**: skills that teach API *vocabulary* (naming patterns, architectural concepts, migration strategies) without providing API *ground truth* (complete method signatures, exhaustive exception hierarchies, valid version numbers) create gaps the model fills with fabrication. The skill provides enough pattern structure for confident generation but insufficient constraint to keep generation within the bounds of real APIs. Grounded tasks are immune to this effect (upgrade-stripe task 04, D-A = 0.000) because the concrete code context supplies the missing constraint.

If this hypothesis holds, the implication is that **more complete API documentation should reduce fabrication**, a prediction that is directly testable by augmenting upgrade-stripe with reference files containing actual

Stripe Python SDK documentation (method signatures, exception class hierarchy, valid API versions) and measuring whether fabrication signals decrease. However, this trades one risk for another: skills with comprehensive reference files are prone to token budget competition and output inflation (the provider-resources pattern, mechanism 3 above). The practical question is whether there exists a middle ground (enough documentation to constrain fabrication without enough to trigger inflation) or whether these failure modes are inherently in tension.

Experimental test: targeted vs. comprehensive reference files. To test the partial knowledge hypothesis directly, we created two synthetic variants of the upgrade-stripe skill with identical SKILL.md files but different reference documentation levels: *targeted* (~2K tokens: version numbers, client initialization, error hierarchies, and webhook signatures across four languages) and *comprehensive* (~8.6K tokens: full SDK documentation including pagination, object expansion, idempotency, retry patterns, and complete webhook handlers).

On the original five evaluation tasks, targeted reference files eliminated the mean B-A degradation entirely (B-A: +0.000 vs. -0.117 baseline) and reduced realistic-condition degradation by 88% (D-A: -0.04 vs. -0.31). Comprehensive references performed worse than baseline on both metrics (B-A: -0.15, D-A: -0.15), with the judge flagging cross-SDK pattern leakage (Python code adopting Node.js-style parameter conventions from the multi-language reference files). The api_idiomaticity dimension showed the clearest dose-response: mean B-A of -0.17 (no refs), -0.33 (targeted), -0.58 (comprehensive), suggesting that more cross-language documentation increases rather than decreases API idiom confusion.

However, the primary risk of reference files was not fabrication but **hybridization**: the model adopted new API vocabulary from references (e.g., `StripeClient`, `v1.customers.create`) but filled procedural gaps from pretrained knowledge of the older API, producing chimeric code with correct class names and incorrect calling conventions. This is a distinct failure mode from the pure fabrication observed without reference files.

Probing evaluation sensitivity. The targeted reference result raised a validity concern: did the targeted references succeed because they addressed genuine knowledge gaps, or because they happened to cover exactly the API surfaces our five tasks were designed to test? To investigate, we added three out-of-band tasks exercising Stripe API surfaces not covered by any reference file: PaymentIntent creation and refund handling (Python), Connect account onboarding (Node.js), and usage-based billing meter events (Ruby). These tasks were designed with knowledge of what the references *don't* cover, creating a controlled test of whether the reference's benefit generalizes.

The results were unambiguous: targeted's advantage did not generalize. On the three out-of-band tasks, targeted showed mean B-A of -0.500 (vs. +0.000 on the original tasks), worse than comprehensive's -0.361. Every out-of-band B-A delta was negative for both variants. The damage concentrated in api_idiomaticity and code_quality, with the judge flagging Python-specific patterns (e.g., `event.identifier`) leaking into Ruby code where `event.id` is correct.

This finding has two implications. First, for the partial knowledge hypothesis: reference files' protective effect is **local to the API surfaces they cover**. Partial ground truth does not create a general grounding effect; it addresses specific fabrication gaps while leaving the broader contamination mechanism (cross-language pattern injection from the SKILL.md) intact. Second, for the behavioral evaluation methodology: because our tasks were designed to exercise specific contamination vectors identified from skill content, the measured degradation represents the effect at maximum sensitivity. The eval answers "can this skill cause this type of interference?" rather than "how much interference does this skill cause across the range of tasks a user would perform?" This is a standard experimental trade-off: targeted probes provide high sensitivity for mechanism detection at the cost of ecological validity for magnitude estimation.

The interaction between novelty, partial knowledge, and fabrication warrants further investigation with larger behavioral samples. Pattern-level guidance without grounding documentation may be worse than either no skill at all (where the model hedges) or a comprehensive skill (where the model is constrained), but confirming this requires testing across a broader range of skills and API domains.

3.4 Company vs. Community: Quality Comparison

A key question motivating this expanded analysis was whether company-published skills demonstrate higher quality than community contributions. The answer is mixed:

Dimension	Company (288)	Community Collections (167)	Anthropic (16)
Pass rate	79.2%	94.0%	87.5%
Avg tokens	6,790	22,900	8,189
Avg info density	0.266	0.174	0.148
Avg specificity	0.585	0.579	0.725
Avg contamination score	0.127	0.091	0.078

Companies produce more **informationally dense** skills (higher code-to-prose ratio) but score lower on **structural compliance** (79.2% vs 94.0% for community collections) and **instruction specificity**. This suggests companies prioritize API reference content over the instructional framing that helps agents use the information effectively.

Anthropic's own skills, while small in number, set a benchmark for instruction specificity (0.725); their skills use strong directive language that leaves less room for agent misinterpretation. Community collections fall between company and Anthropic skills on most dimensions.

3.5 Metric Correlations

Notable correlations: - Token count shows weak positive correlation with error count, suggesting larger skills tend to have more structural issues - Information density and code block ratio are strongly correlated (by construction) - Risk score correlates with warnings, as structurally complex skills tend to have broader technology scopes

3.5.1 LLM Dimension Correlations

The six LLM judge dimensions form two distinct clusters. The five *craft quality* dimensions (clarity, actionability, token efficiency, scope discipline, and directive precision) intercorrelate at r = 0.22–0.61, forming a coherent quality factor. Novelty stands apart, correlating with craft dimensions at only r = 0.04–0.39. This two-factor structure confirms that novelty captures a genuinely independent quality signal: a skill's writing craft and its information novelty are largely orthogonal.

Cross-correlating LLM dimensions with heuristic metrics reveals that information density is the best heuristic predictor of LLM-judged quality (r ≈ 0.50 with actionability), validating our heuristic as a useful proxy. Token efficiency anti-correlates with word count (r ≈ −0.45); longer skills are penalized by the judge, consistent with the view that verbosity degrades agent performance. Contamination score and novelty are essentially
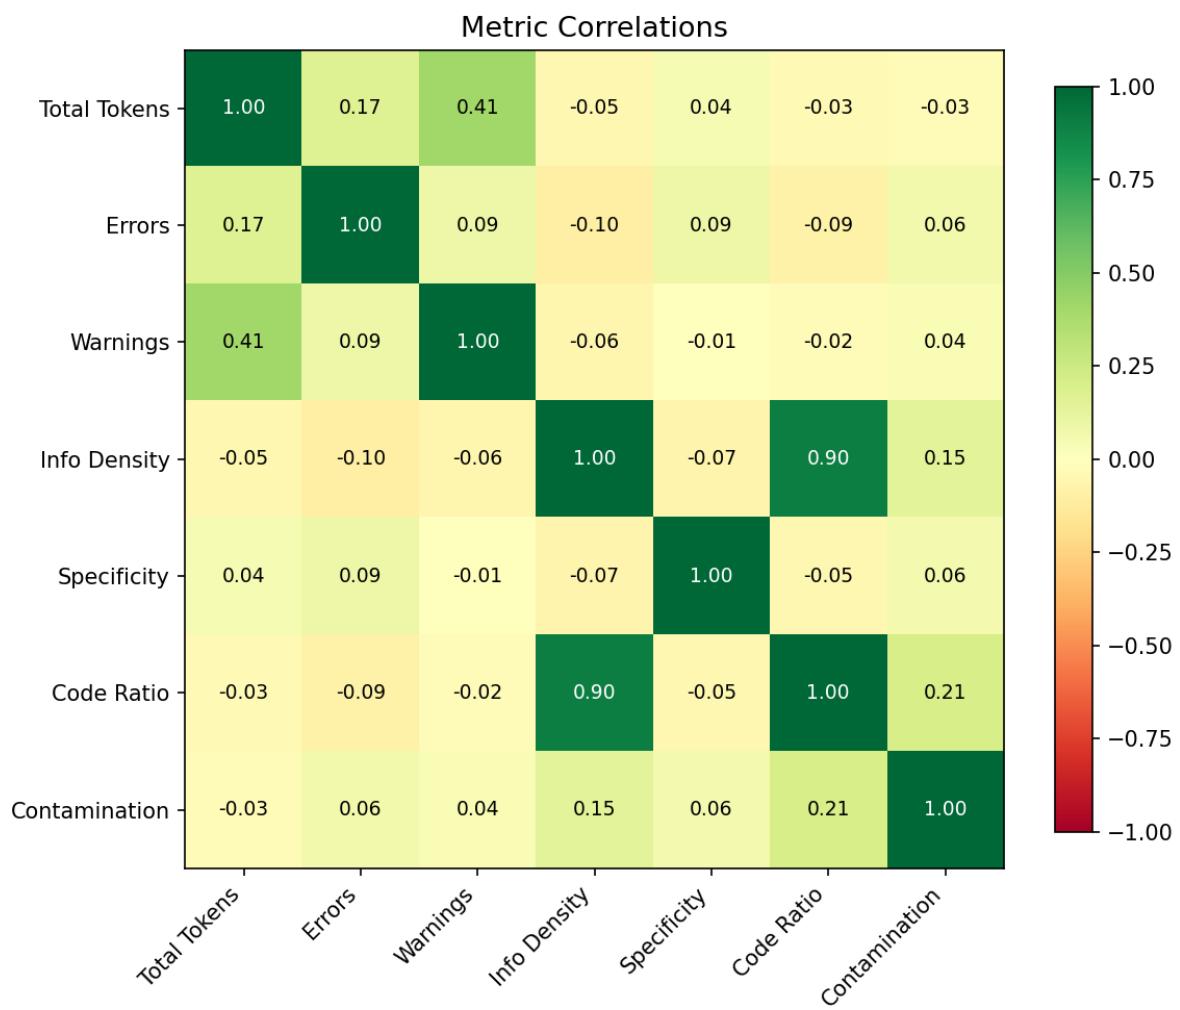
Figure 14: Correlation between key metrics
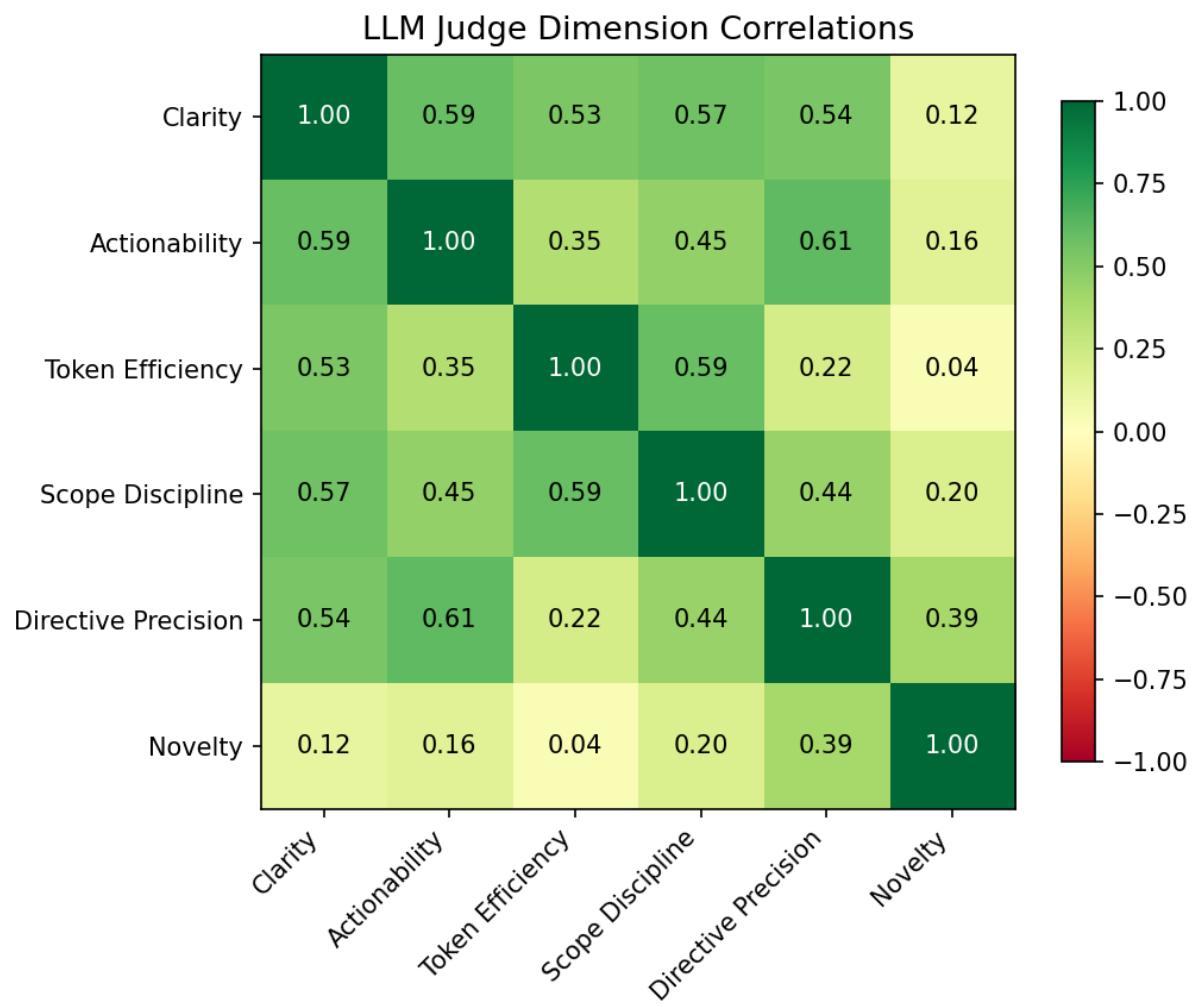
Figure 15: Correlations between LLM judge dimensions
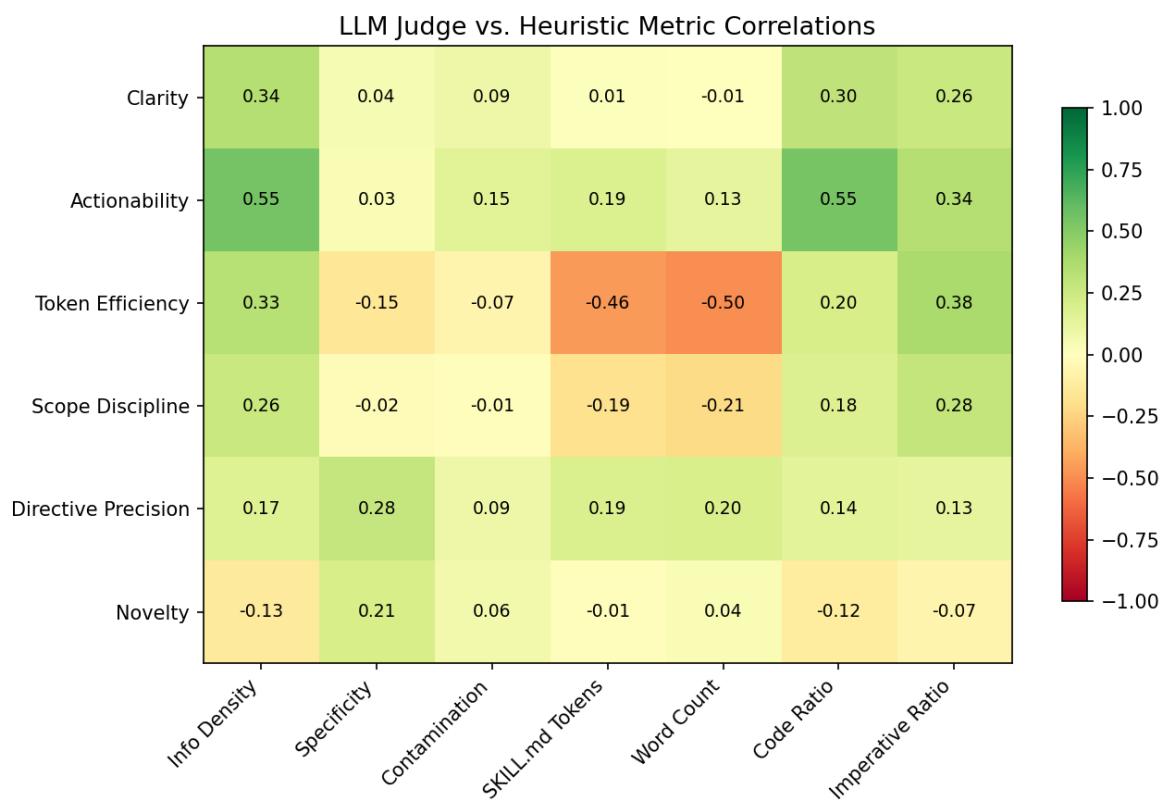
Figure 16: LLM judge vs. heuristic metric correlations

uncorrelated (r ≈ 0.07), confirming that these metrics are complementary rather than redundant: contamination measures *structural* risk from mixed-language content, while novelty measures *informational* value beyond training data.

3.6 Reference File Analysis

Beyond the primary `SKILL.md` file, the agentskills.io specification allows skills to include reference files, supplementary documents placed in a `references/` directory. These files provide additional context (API documentation, code examples, configuration templates) that agents load alongside the skill instructions. Our analysis reveals that reference files represent an underexamined dimension of skill quality.

3.6.1 Prevalence and Scale

Of 673 skills analyzed, **412 (61%) include reference files**, totaling 1,877 files across the dataset. Reference file usage varies by source: company-published skills and community collections are the heaviest users, while methodology-focused skills (K-Dense) and security skills rarely include references.

3.6.2 Token Budget Impact

Reference files have a dramatic impact on context window consumption. Among skills with references:

- **81% have more reference tokens than SKILL.md tokens**, meaning the references outweigh the primary instruction file
- The median reference token count is 4,729, but the distribution has a heavy tail: p99 reaches 44k tokens
- **4 skills have reference totals exceeding 50,000 tokens**, prime candidates for context window degradation. Controlled experiments show 13.9–85% performance loss as input length increases, even when models can perfectly retrieve all evidence (Du et al. 2025)
- Extreme outliers exist: vueuse-functions (153k reference tokens, 18x the SKILL.md) and security-best-practices (91k reference tokens, 57x the SKILL.md)

In practice, agent platforms typically operate within context windows of 100k–200k tokens, though research suggests that effective context lengths are often less than half the advertised training length due to undertrained positional encodings (An et al. 2025). A single oversized reference file can consume a significant fraction of this budget, potentially crowding out the user's code context and degrading agent performance, an effect documented across 18 state-of-the-art models, where performance degrades non-uniformly as context grows (Hong et al. 2025; Yoran et al. 2024). Skill authors should be mindful that reference files are not "free"; they compete directly with the user's code for context window space (Salim et al. 2026).

3.6.3 Content Quality

Reference files tend to have **higher information density** than their corresponding SKILL.md files. This is expected: references are typically code-heavy (API examples, configuration templates, type definitions) rather than prose-heavy instruction documents. The higher density reflects their role as reference material rather than instructional content.

LLM-as-judge scoring confirms and quantifies this pattern. Across 412 skills with both SKILL.md and reference scores, reference files outscore their parent SKILL.md on every shared dimension: +0.48 overall, +0.66 on clarity, +0.62 on token efficiency, and +0.19 on novelty. References are more concise and information-dense (focused code examples and API documentation versus prose instructions), which the LLM judge
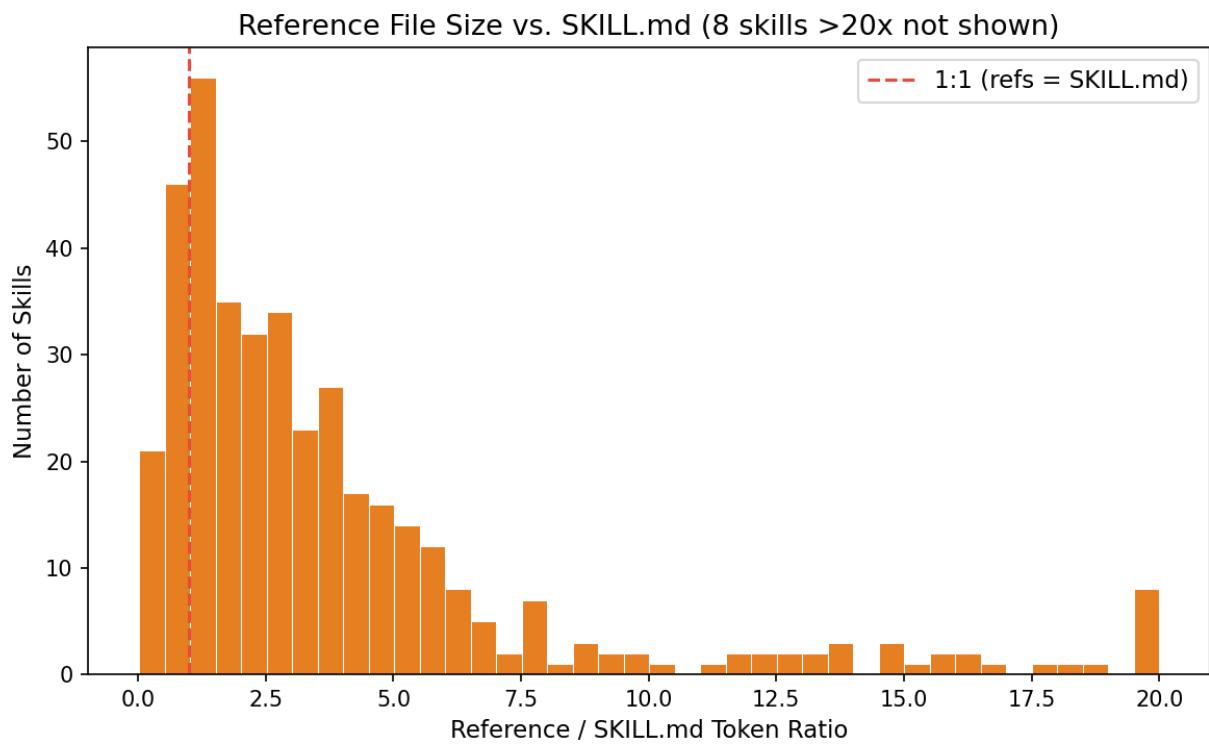
Figure 17: Reference file size relative to SKILL.md
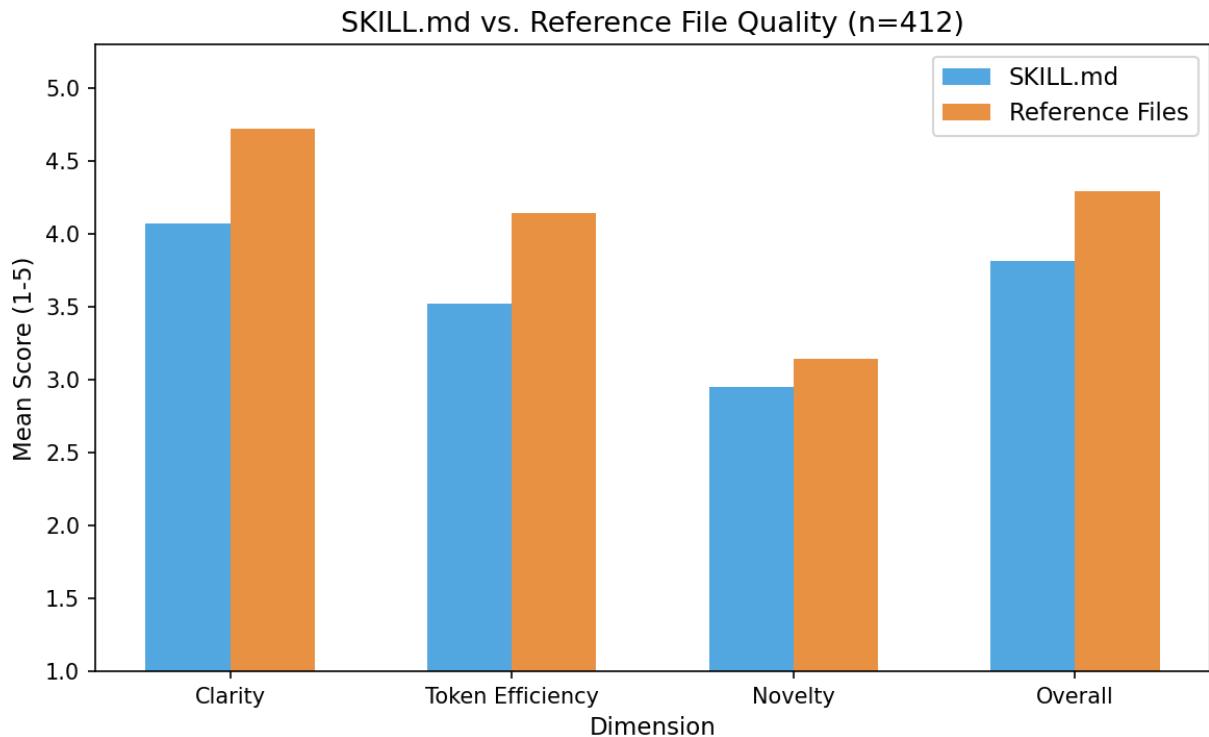
Figure 18: SKILL.md vs. reference file quality comparison

rewards. The novelty gap is negligible, suggesting that while references contain more *efficiently presented* information, the *uniqueness* of that information relative to training data is comparable to SKILL.md content.

3.6.4 Hidden Contamination in Reference Files

Reference files introduce a distinct contamination vector. Because references often contain code examples for specific language SDKs, they may interfere with the agent's code generation when the user is working in a different language, a concern grounded in research showing that in-context code examples bias generated output toward the patterns present in those examples (Li et al. 2025; Ali et al. 2024). Our analysis found that reference contamination patterns differ from SKILL.md contamination:

- Reference files are more likely to contain **multiple programming languages** within a single skill, especially for platform SDK skills that provide examples across Python, TypeScript, Java, and .NET
- Skills with contaminated reference files sometimes have clean SKILL.md files; the contamination is hidden in the supplementary material

Among the 412 skills with reference files, contamination scores diverge between the SKILL.md and references in the majority of cases:

- **32.3%** have higher contamination in references than in SKILL.md
- **66.3%** have higher contamination in SKILL.md than in references
- **1.5%** have equal scores

The most concerning pattern is what we term **hidden contamination**: skills where the SKILL.md file scores as low-risk but the reference files carry medium or high contamination. We identified **66 skills** (16.0% of skills with references) exhibiting hidden contamination: 12 with high-risk references and 54 with medium-risk references.

Community collections account for the majority of hidden contamination cases (31 of 66), which is expected given their heavy reliance on reference files (88% adoption rate, averaging 6.9 files per skill). Company-published skills contribute 23 cases, and vertical/domain-specific skills contribute 6.

Illustrative examples of hidden contamination:

- **neon-postgres** (company): SKILL.md contamination 0.00 (low), reference contamination 0.83 (high). The SKILL.md provides clean PostgreSQL-focused instructions, but reference files include examples spanning JavaScript, Python, CSS, TSX, and Bash, with the JavaScript/Python application-language mixing driving the high reference score.
- **react-native-best-practices** (company): SKILL.md contamination 0.07 (low), reference contamination 1.0 (high). Reference files contain code in 14 different languages including Kotlin, Swift, Objective-C, Groovy, and C++, a mix of application and mobile language categories that produces the maximum contamination score.
- **alphafold-database** (community collection): SKILL.md contamination 0.03 (low), reference contamination 0.21 (medium). Reference files mix Python with Bash and configuration languages for the protein structure prediction pipeline. The similarity-weighted scoring appropriately reduces this from high to medium, as the Python/bash mixing is primarily an auxiliary mismatch.
- **clinical-decision-support** (community collection): SKILL.md contamination 0.00 (low), reference contamination 0.41 (medium). A perfectly clean instruction file is paired with references containing mixed-language code examples.

The consequence for quality assessment is direct: **any contamination analysis that examines only the SKILL.md file will miss 66 cases**, a 30% undercount relative to the 223 total skills with medium or high
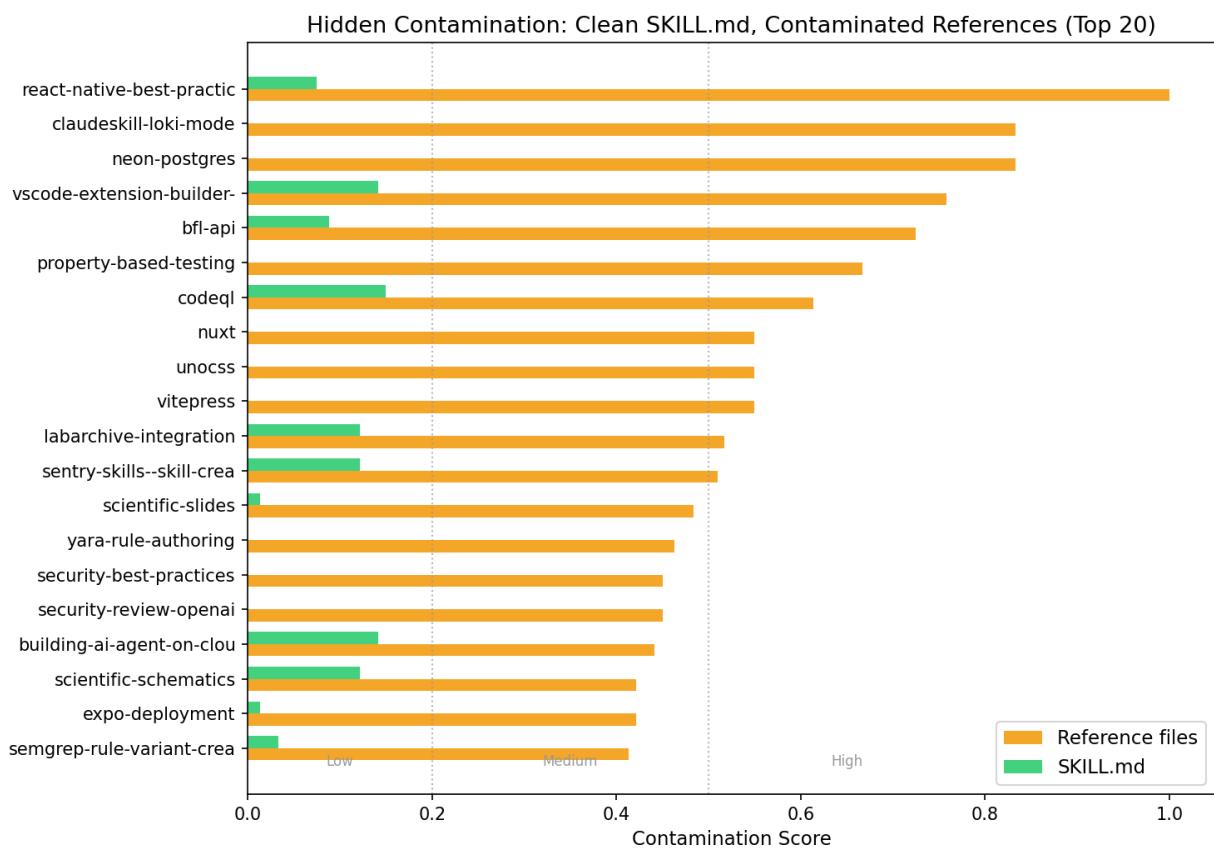
Figure 19: Hidden contamination: clean SKILL.md with contaminated references

contamination across either their SKILL.md or references. Skill validation tools and quality gates should evaluate reference files alongside the primary instruction file.

3.6.5 Language Distribution

Most Common Languages in Reference Files (Top 15)
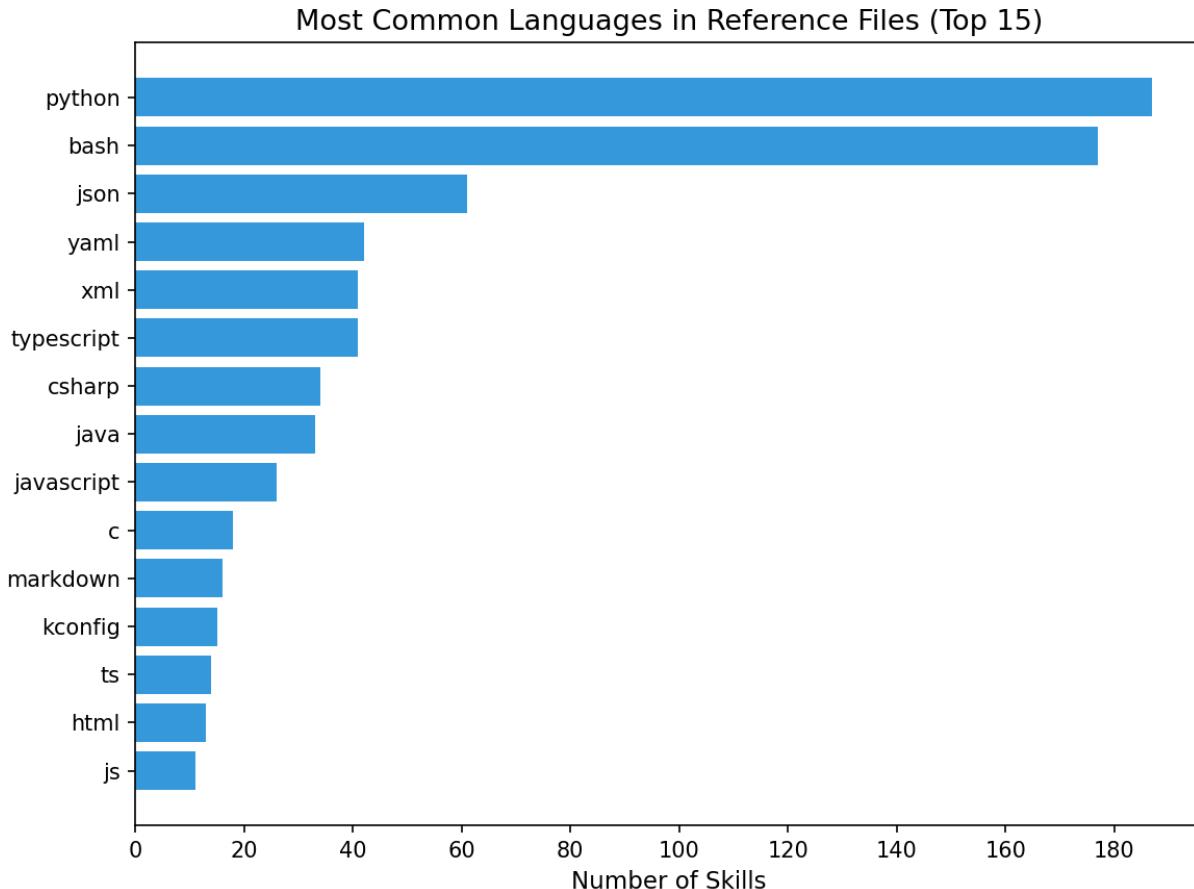
Figure 20: Most common languages in reference files

The language distribution across reference files reveals how authors use references in practice. The most common languages reflect the ecosystem's emphasis on web development and cloud platform integration. Shell scripts, configuration languages (YAML, HCL, TOML), and type definition files appear frequently, confirming that references serve as practical implementation guides rather than conceptual documentation.

4 Recommendations for Skill Authors

Based on our findings, we organize recommendations into three groups: practices where our data validates existing guidance from the specification maintainer (Anthropic), practices where our data suggests existing guidance needs more detail, and practices that address dimensions not covered by existing guidance.

4.1 What Existing Guidance Gets Right

Anthropic's `skill-creator` skill (Anthropic 2025) provides sound authoring principles. Our analysis of 673 skills provides ecosystem-wide evidence for how well these principles are followed, and quantifies the consequences when they are not.

1. **Prioritize novel content over restating common knowledge**: Anthropic's skill-creator advises authors to "only add context Claude doesn't already have" and to challenge each piece of information with "Does Claude really need this explanation?" Our data shows the ecosystem largely ignores this guidance: novelty is the weakest dimension across all 673 skills (mean 3.12/5) and the key differentiator between top and bottom skills. Community collection skills score lowest (2.83), precisely because they repackage information already in the model's training data. Skills wrapping well-known libraries should focus on non-obvious gotchas, internal conventions, and project-specific patterns rather than restating standard documentation.

2. **Minimize token usage**: The skill-creator correctly identifies the context window as "a public good" shared with system prompts, conversation history, and user code. Our data quantifies the problem: the median effective skill uses ~5,200 tokens, yet 52% of all tokens ecosystem-wide are nonstandard files (LICENSE texts, build artifacts, XML schemas) that provide no instructional value. Skills exceeding 50,000 tokens likely include material that should be external references or removed entirely.

3. **Delegate depth to reference files**: Anthropic provides detailed progressive disclosure patterns (domain-specific organization, conditional detail loading, and high-level guides with references). Our LLM-as-judge data validates this approach: reference files consistently outscore SKILL.md (+0.48 overall, +0.62 on token efficiency). Keep SKILL.md at 100–300 lines as a concise action guide and delegate detailed examples, API docs, and extended patterns to `references/`. Top-scoring skills use this pattern; bottom-scoring skills either dump everything into SKILL.md or contain nothing but pointers to references.

4. **Audit nonstandard files**: The skill-creator warns against extraneous files (README.md, CHANGELOG.md, etc.). Our analysis reveals the scale of non-compliance: 185 skills (27.5%) contain nonstandard files, inflating mean token counts by 108%. Check that your skill directory contains only `SKILL.md`, `references/`, and `assets/`. Move license text to a reference or remove it; relocate build artifacts and schemas outside the skill directory.

5. **Validate before publishing**: The specification provides validation tools and the skill-creator includes a packaging step that validates automatically. Despite this, 22.0% of published skills fail structural validation, including 21% of company-published skills. Run `skill-validator` on your skill and fix all errors before publishing.

4.2 Where Existing Guidance Needs More Detail

These recommendations build on principles present in Anthropic's guidance but add data-driven specificity not found in the existing documentation.

6. **Use strong directives strategically, not universally**: The skill-creator advises using "imperative/infinitive form." Our analysis refines this: top-scoring skills use strong markers (MUST, ALWAYS, NEVER) at high-consequence decision points (security boundaries, common failure modes, prerequisite checks) rather than scattered throughout as emphasis for routine instructions. Company skills in particular tend toward advisory language ("consider", "might") where directive language would be more effective, but indiscriminate use of strong directives reduces their signal

value.

7. **Write a rich frontmatter description that naturally incorporates keywords**: The specification requires a description field and says it "should include specific keywords that help agents identify relevant tasks." The skill-creator emphasizes it as "the primary triggering mechanism." Our analysis adds a specific recommendation: 50–200 words of natural prose that weaves in relevant keywords organically. Avoid bare keyword lists (`MongoDB, Atlas, Vector Search, embeddings`); this comma-separated keyword pattern appears in 100 skills across the ecosystem. Explicit trigger phrase lists (`Triggers: "term1", "term2"`) are less harmful when accompanied by substantive prose but still less readable than descriptions that weave keywords naturally into sentences. The skill-creator also correctly notes that "When to Use" information belongs in the description rather than the body, since the body is only loaded after the skill triggers. Body-level scope sections (see recommendation #12) serve a different purpose: helping the agent apply the skill correctly once activated.

4.3 New Recommendations

These practices address dimensions not covered by Anthropic's existing guidance, either entirely new concerns identified by our analysis or structural patterns derived from ecosystem-wide scoring.

8. **Scope skills tightly to avoid content interference**: Skills covering multi-interface tools should target a specific language SDK. A "MongoDB for Node.js" skill is safer than a generic "MongoDB" skill. Our exploratory behavioral evaluation (n = 19) found no correlation between structural contamination scores and behavioral degradation (r = 0.077), and case studies suggest content specificity may matter more: framework-specific patterns interfered with unrelated tasks, while language-portable patterns (e.g., security vulnerability patterns across languages) did not. The practical guidance remains the same: tight scoping reduces both language confusion risk (Moumoula et al. 2026) and the content-specific interference our behavioral eval identifies.

9. **Label code blocks explicitly**: Always specify the language in fenced code blocks (```` ```javascript ```` rather than ```` ``` ````). This is good practice for readability and parseability, and research shows explicit language keywords mitigate Programming Language Confusion within a single context (Moumoula et al. 2026). However, in our 19-skill behavioral sample, we observed no protective effect of language labeling on cross-task interference: skills with 100% label rates showed worse mean degradation (-0.105) than partially-labeled skills (-0.052), though this comparison is confounded by other skill differences and should not be interpreted causally. Labels address within-context language disambiguation but may not address the content-specific interference mechanisms (template propagation, textual frame leakage) we observed.

10. **Separate language-specific examples**: If a skill must cover multiple languages, use clearly delineated sections with explicit context-switching markers. Consider publishing separate skills per language SDK. Our Gemini API case study illustrates how four subtly different API shapes for the same operation create the conditions for syntactically plausible but semantically incorrect code generation.

11. **Favor minimal patterns over comprehensive examples in platform skills**: Enterprise platform skills (SDK guides, provider development frameworks, infrastructure tools) are especially prone to output inflation, architectural pattern bleed, and API hallucination. The **provider-resources** skill (HashiCorp) teaches detailed Go Terraform provider conventions through comprehensive code examples; when loaded, the model generates 2× longer implementations that hit token limits and get truncated, and transfers the skill's Go package architecture (separate models, clients, helpers) into Python outputs where a simpler structure would be idiomatic. Similarly, **react-native-best-practices**

(Callstack) causes the model to allocate 2.6× more tokens to framework-specific commentary at the expense of code completeness. The **upgrade-stripe** skill (Stripe) demonstrates a distinct enterprise risk: its detailed API versioning and migration content causes the model to fabricate plausible but incorrect SDK calls (e.g., nonexistent version strings like `stripe-go/v81`, incorrect method signatures, and hallucinated Go error constants), and this fabrication *intensifies* under realistic agentic context (D-A = -0.383 vs. B-A = -0.117), the opposite of the typical mitigation pattern. The mechanism holds up: we validated it by controlling for a codebase snippet confound that initially inflated the effect, and the amplification persists across Python, Go, JavaScript, and Ruby tasks. For enterprise users, this means agent-generated code becomes more verbose and complex (increasing maintenance burden), non-idiomatic architectural patterns bleed across language boundaries (creating technical debt), and skills teaching API surface knowledge can make the model *more confidently wrong*, turning uncertain hallucinations into plausible fabrications that are harder to catch in code review. Platform skill authors should prefer concise pattern summaries over complete implementation examples, and should favor showing the minimal correct pattern for each operation rather than a comprehensive reference implementation. Skills that teach "here is the complete way to build X" are more prone to these effects than skills that teach "here are the key patterns to follow when building X."

12. **Use structured formats over prose**: Top-scoring skills use tables as their primary information vehicle (90% include tables vs. 40% of bottom skills) and maintain roughly a 4:1 ratio of structured content (tables, lists, code) to prose. Quick reference tables near the top of the file, anti-pattern tables, and decision matrices make skills parseable by agents without requiring sequential reading. Every bottom-10 skill in our analysis is prose-heavy. Anthropic's own skill-creator guidance is itself prose-heavy and does not mention tables as a structural pattern.

13. **Include explicit anti-patterns and negative scope**: Document what NOT to do and why (anti-patterns section), and define what is out of scope (negative scope gate). Of the top 20 skills, 60% include an explicit anti-patterns section and 50% include "When NOT to Use" guidance; both are absent from every bottom-10 skill. Negative guidance is not mentioned in Anthropic's existing authoring documentation.

14. **Follow the empirical template structure**: Our analysis of top-scoring skills reveals a convergent architecture: scope gate → quick reference table → core workflow → domain-specific sections → anti-patterns → troubleshooting → references. We provide a concrete proposed template derived from these patterns (see companion repository). Of the top 20 skills, 80% include a quick reference section near the top; this pattern is absent from every bottom-10 skill and from Anthropic's existing template and guidance.

15. **Validate reference files for contamination**: A clean SKILL.md does not guarantee a clean skill. We identified 66 skills with hidden contamination in reference files, a 30% undercount relative to total contaminated skills if only SKILL.md is analyzed. Run contamination analysis on the full skill directory, not just the instruction file.

16. **Evaluate whether low-novelty skills add value**: Our exploratory behavioral evaluation (n = 19) observed a suggestive pattern where novelty correlates with the magnitude of behavioral effects (r = +0.267 for |B-A|, p ~ 0.27), with low-novelty skills showing smaller effects overall, as if the model largely ignores content it already knows. If this pattern holds, it would partially mitigate the "net negative" concern for the 51 low-novelty, medium-to-high contamination skills we identified. However, the risk is not zero, and a skill that provides no novel information while consuming context window budget is still inefficient. Before publishing, authors should consider whether the content adds value beyond what the LLM already knows.

5 Recommendations for Spec Maintainers

The agentskills.io specification ("Agent Skills Specification" 2025) provides structural requirements (directory layout, frontmatter fields, file organization), and Anthropic's `skill-creator` skill (Anthropic 2025) offers sound authoring principles (context window efficiency, progressive disclosure, novelty over repetition). Our recommendations focus on gaps: areas where the specification and existing guidance are silent, and where our ecosystem-wide data suggests intervention would have the greatest impact.

1. **Add a `languages` frontmatter field**: Skills should explicitly declare which programming languages they target. This enables agents to filter skills by context and would help mitigate cross-contamination. Neither the specification nor the skill-creator currently addresses language scoping.

2. **Define quality tiers with separate craft and novelty axes**: Introduce a quality score based on structural compliance, content metrics, contamination risk, and LLM-judged dimensions. Our analysis reveals a two-factor structure: five craft dimensions (clarity, actionability, token efficiency, scope discipline, directive precision) intercorrelate at r = 0.22–0.61, while novelty is largely independent (r = 0.04–0.39). Quality tiers should reflect both axes, since a skill can be well-crafted but unoriginal, or novel but poorly written, and the improvement path is different for each.

3. **Require code block language annotations**: Make unlabeled code blocks a validation error, not just a warning. Language annotations improve readability, enable syntax highlighting, and aid programmatic analysis. Research shows explicit language keywords mitigate Programming Language Confusion within a single context (Moumoula et al. 2026), though our exploratory behavioral evaluation (n = 19) found no protective effect against the cross-task interference mechanisms observed in our sample (see Behavioral Evaluation Results).

4. **Provide multi-language skill guidelines**: Neither the specification nor the skill-creator addresses skills that necessarily cover multiple languages (CI/CD, infrastructure, cross-platform tools). With 10 high-risk and 147 medium-risk skills in our sample, this is an important gap.

5. **Add content interference assessment**: Include cross-contamination detection in the specification's recommended validation pipeline. Our exploratory behavioral evaluation found no correlation between structural contamination scores and behavioral degradation (r = 0.077, n = 19), suggesting that validation should go beyond language-mixing heuristics to flag content-specific risks: invalid syntax in output templates, framework-specific patterns that may bleed into unrelated contexts, and suggestive API patterns that could trigger hallucination. Critically, this assessment should cover reference files as well as SKILL.md, since our analysis found 66 cases of hidden contamination visible only in references.

6. **Engage company publishers on novelty, community authors on craft**: Company-published skills lead on craft quality (scope discipline 4.74, clarity 4.13, token efficiency 3.86) but rank #5 of 8 sources on novelty (3.12); they are polished documentation for tools the LLM already knows. Community collections show the inverse: weakest on novelty (2.83) but reasonably actionable. The improvement path differs by source: company publishers should focus on non-obvious integration patterns and gotchas rather than restating their public API docs, while community authors would benefit most from structural templates and editorial guidelines to improve craft. The skill-creator's advice to "only add context Claude doesn't already have" is sound but insufficiently emphasized; our data shows it is the most-ignored principle in the ecosystem.

7. **Warn on or penalize nonstandard files**: The skill-creator warns against extraneous documentation files, but the specification itself does not warn against or penalize nonstandard files at the validation

level. Currently, 52% of all tokens in the ecosystem come from nonstandard files, including 75.1% of tokens in Anthropic's own skills. A validation warning or error for unexpected root-level files would significantly reduce context window waste. At minimum, agent platforms should consider filtering out nonstandard files when loading skills.

8. **Set token budget guidelines**: The specification recommends "< 5000 tokens" for SKILL.md and the skill-creator says "under 500 lines," but neither provides guidelines for reference files or total skill size. Our analysis shows that the median effective skill is ~4,000 tokens, yet 17 skills exceed 50% of a 128k context window. Guidelines for individual reference files and total skill token budget would help authors understand the practical limits of context window consumption.

9. **Publish a substantive SKILL.md template with structural patterns**: The current official template provides three lines of scaffolding. The skill-creator offers good *principles* (conciseness, progressive disclosure, novelty) but does not prescribe a specific *structure*; it says "write instructions" without specifying which sections to include or in what order. Our analysis of top-scoring skills reveals structural patterns absent from existing guidance that consistently distinguish high-quality skills: quick reference tables (80% of top-20, 0% of bottom-10), anti-pattern documentation (60% of top-20, 0% of bottom-10), negative scope gates, and a convergent section architecture. A template encoding these patterns, with inline guidance explaining *why* each section matters, would operationalize the skill-creator's principles into concrete structure that authors can follow. We provide a proposed template derived from these empirical findings as a companion to this paper.

6   Limitations and Future Work

Limitations:

- Our content metrics (information density, instruction specificity) are heuristic and may not capture all aspects of skill quality
- Cross-contamination risk scoring measures structural indicators (multi-language content, multi-interface tools, scope breadth) using keyword matching rather than semantic analysis. Our exploratory behavioral evaluation of 19 skills found no correlation between these structural scores and behavioral degradation ($r = 0.077$, $n = 19$), and the interference mechanisms observed in case studies were content-specific rather than language-mixing artifacts. The structural scores remain useful as indicators of multi-language complexity but should not be interpreted as behavioral risk predictions
- The scoring weights for language-type mismatches (application-to-application: 1.0, application-to-auxiliary: 0.25, auxiliary-to-auxiliary: 0.1) are informed by the research direction rather than empirically calibrated. Our exploratory behavioral evaluation ($n = 19$) suggests that the weighting scheme may capture the wrong dimension: the case studies indicate content specificity and task mismatch may matter more than language similarity for predicting interference
- The behavioral evaluation covers 19 skills (a 2.8% sample) with 5 tasks each, evaluated at temperature 0.3 with 3 runs per condition. While this provides statistical power for individual skill effects, the sample is too small for robust cross-skill correlations. All correlations reported (e.g., novelty vs. |B-A|, contamination vs. B-A) are exploratory and should be interpreted cautiously
- The behavioral evaluation's task design is inherently targeted: tasks and anti-patterns were crafted with knowledge of each skill's content to maximize sensitivity to the specific contamination mechanisms each skill could introduce. This means the eval functions as a hypothesis confirmation test: it can detect the mechanisms it was designed to detect, but may miss degradation types that were not anticipated (e.g., subtle architectural over-engineering, increased verbosity, or hallucination of APIs not

present in the skill). Our out-of-band experiment with upgrade-stripe confirms this limitation: reference file improvements that eliminated degradation on the five original tasks (B-A: +0.000) produced worse-than-baseline results on three tasks exercising uncovered API surfaces (B-A: -0.500). The LLM judge provides a partially independent signal (it scores on generic dimensions without knowledge of anti-patterns), but its prompt framing ("evaluating code for signs of cross-language contamination") primes it toward the same failure mode the eval was designed to detect. The negative controls (fastapi-router-py and doc-coauthoring, both with contamination score 0.00) provide some protection against systematic inflation: if the methodology itself generated false positives, the controls would show it. All behavioral findings should be interpreted as mechanism-detection results at maximum sensitivity, not as estimates of real-world degradation magnitude

- The behavioral eval uses a fixed 4,096-token output ceiling. 63.9% of baseline runs and 62.0% of with-skill runs produce assessments mentioning truncation, making this a pervasive confound rather than an edge case. However, a floor-effect analysis finds that truncation does not mask stronger degradation signal: tasks where both conditions are truncated show a mean B-A of -0.140 (slightly *worse* than the -0.080 overall mean, likely because skill context consumes output token budget), while tasks where neither condition is truncated show a mean B-A of only -0.046, essentially noise. Score ranges are compressed in truncated tasks (1.25 vs. 3.08 for clean tasks) and variance is mildly reduced (~14%), but the structural contamination correlation remains near zero in the clean subset (r = -0.065 vs. r = 0.077 overall). Truncation slightly inflates apparent degradation through token budget competition rather than masking a stronger underlying signal

- For skills with large reference directories, the behavioral eval loads 2–3 reference files per task rather than all references. This selective loading was necessary to avoid measurement artifacts (loading all references produced near-identical outputs, effectively reducing n to 1), but means the eval tests a curated subset of each skill's content rather than the full skill as a user would encounter it

- Condition D (realistic context) injects language-specific codebase snippets as simulated conversation history. We identified one confound: the default Python snippet uses async patterns (FastAPI + Async-Session) that caused the LLM judge to flag async/sync confusion as skill-induced contamination in upgrade-stripe's Python task. Re-running with a synchronous snippet reduced that task's D-A delta by ~70%. Other language snippets may contain similar confounds; we corrected the identified case but did not systematically audit all 12 language snippets

- Skills that contain embedded AI-directed prompts or evaluation rubrics pose a distinct quality risk beyond token waste or contamination: they can act as unintentional prompt injections when loaded into an agent's context window. We observed this directly during LLM-as-judge scoring, where three skills consistently caused the judge model to abandon its evaluation task. A meta-skill containing realistic pressure scenarios (e.g., "Choose A, B, or C. Be honest." and "Make agent believe it's real work, not a quiz") caused the judge to role-play the embedded scenarios rather than evaluate the document. A diagram generation skill with an embedded quality rubric (scoring criteria like "Scientific Accuracy (0-2 points)" and example output showing "SCORE: 8.0") caused the judge to adopt that rubric format instead of its own. Reference files with code-heavy content (particularly C# examples containing curly-brace initializers, inline JSON schemas, and string interpolation) consistently corrupted the judge's JSON output formatting, producing syntactically invalid responses. While an agent performing a user-directed task has a richer context that provides more insulation than a single-purpose judge, the risk is not zero, particularly for skills containing embedded evaluation criteria, example dialogues, or dense code with JSON-like syntax. These three failure modes (prompt hijacking via embedded scenarios, format hijacking via embedded rubrics, and output corruption via code-heavy content) represent categories of skill content that structural validation cannot detect but that skill authors should be aware of

- The dataset represents a point-in-time snapshot; the ecosystem evolves rapidly

- Some skill name collisions occur across sources within the same category

Future work:

- **Broader behavioral coverage**: Our exploratory behavioral evaluation of 19 skills found no correlation between structural contamination scores and behavioral degradation, and identified six distinct content interference mechanisms. Expanding behavioral testing to a larger sample, particularly targeting the 51 "net negative" skills with low novelty and high contamination, would test the robustness of the zero correlation, strengthen the suggestive cross-skill patterns (e.g., novelty amplification), and determine whether low-novelty skills truly produce smaller effects than high-novelty skills
- **Calibrating structural risk scores**: The disconnect between structural contamination scores and behavioral impact ($r = 0.077$) suggests that the scoring heuristic should be revised. Future work could develop content-specific risk indicators based on the mechanisms identified in our behavioral evaluation: template defect density, framework-specificity measures, and API pattern diversity
- **Extending the partial knowledge test**: Our experimental test of the partial knowledge hypothesis (two synthetic upgrade-stripe variants) confirmed that reference files reduce fabrication on covered API surfaces but introduced a new failure mode (hybridization) and showed that the protective effect does not generalize beyond covered surfaces. Expanding this experiment to additional skills with different contamination profiles would test whether these patterns are specific to multi-language SDK skills or represent a general dynamic of reference file interaction. The out-of-band methodology we developed (testing skill modifications against tasks outside the modification's coverage) could also be applied to validate other behavioral eval findings
- **Longitudinal analysis**: Track skill quality trends as the ecosystem matures and companies update their skills
- **Skill composition analysis**: Study how multiple active skills interact and potentially conflict. Our behavioral eval tests skills in isolation; in practice, agents may load multiple skills simultaneously, creating interaction effects that single-skill testing cannot detect
- **Expanded coverage**: Our ecosystem survey (Appendix A) identifies 800+ additional skills; analyzing the full set would strengthen statistical power

7 Conclusion

The Agent Skills ecosystem is young and growing rapidly. Our analysis of 673 skills from 41 repositories reveals meaningful variation in structural compliance (78.0% pass rate), content quality, cross-contamination risk (10 high-risk skills), and context window efficiency (52% token waste). Company-published skills (from Microsoft, OpenAI, Stripe, and others) have a *lower* structural compliance rate (79.2%) than community collections (94.0%), inverting the assumption that official sources produce higher-quality skills.

Three findings merit highlighting. First, the context window waste problem: over half of all tokens loaded from skills are nonstandard files (LICENSE texts, build artifacts, XML schemas, benchmark data) that provide no instructional value to the agent. This is both the largest quality issue by magnitude and the easiest to fix. Second, hidden contamination in reference files: 66 skills appear clean when only the SKILL.md is analyzed but carry medium or high contamination risk in their reference files, meaning contamination assessments that ignore references undercount risk by 30%.

Third, our exploratory behavioral evaluation of 19 skills found **no correlation between structural contamination scores and behavioral degradation** ($r = 0.077$, n = 19). While the sample is too small for definitive cross-skill conclusions, the individual case studies are instructive: a skill with near-zero structural risk (react-native-best-practices, score 0.07) produced the largest behavioral degradation (B-A = -0.383), while a skill

with 12 application languages and 100% code block labeling (sharp-edges, score 0.62) showed minimal effect (-0.083). The interference mechanisms we identified (template propagation, textual frame leakage, API hallucination, and token budget competition) are content-specific rather than language-mixing artifacts, suggesting that the field's focus on multi-language mixing as the primary contamination vector, grounded in PLC research (Moumoula et al. 2026), may capture only one mechanism among several.

Preliminary evidence from this sample also suggests that realistic agentic context (a system preamble plus conversation history) substantially attenuates skill-only degradation (mean mitigation ratio -74.5%), though this estimate is imprecise given the small deltas involved. If this pattern holds at scale, it is reassuring: skills are not loaded in isolation, and the rich context of an agentic workflow may anchor the model's behavior against most interference.

LLM-as-judge scoring across all 673 skills reveals that novelty, the degree to which a skill provides information beyond training data, is the key quality differentiator, correlating only weakly (r = 0.04–0.39) with craft dimensions like clarity and conciseness. This two-factor structure (craft vs. novelty) has practical implications: the improvement path differs by source category, with company publishers needing more novel content and community authors needing better craft. The behavioral evaluation provides suggestive evidence that novelty may also relate to behavioral effects (r = +0.267 for |B-A|, n = 19, not significant), but confirming this requires larger samples.

Quality standards, validation tooling, and authoring guidelines can address these issues. The behavioral case studies suggest that validation should expand from structural language-mixing heuristics toward content-specific checks: template syntax validation, framework-specificity assessment, and reference file scope analysis. Combined with the security vulnerabilities identified by Liu et al. (Liu et al. 2026) (26.1% of skills containing prompt injection, data exfiltration, or privilege escalation risks), the picture is clear: the ecosystem faces quality and safety challenges on multiple fronts. As skills become a core part of the AI development workflow, the community benefits from treating them as first-class software artifacts deserving of the same quality discipline we apply to libraries and APIs.

7.1 AI Assistance Disclosure

The analysis pipeline (data collection, validation, statistical computation, and behavioral evaluation infrastructure) was developed with AI assistance from Claude (Anthropic). The LLM-as-judge scoring was performed by Claude Sonnet. The author designed the study, curated the dataset, interpreted all findings, and wrote the paper; AI tools were used for implementation support, statistical analysis, and iterative refinement of methodology.

Appendix A: Ecosystem Survey

Beyond the 673 skills we analyzed in depth, we conducted a broad survey of the Agent Skills ecosystem to estimate the total scope and identify patterns in adoption. This appendix documents our findings.

7.2 Scale of the Ecosystem

As of February 2026, we identified **120+ repositories** containing Agent Skills, with an estimated **1,400+ individual skills** across the ecosystem. The agentskills.io specification is supported by **27+ agent platforms**.

Our analyzed sample of 673 skills represents approximately 48% of the estimated total. The quality patterns we observe, particularly the 22.0% structural failure rate and 1.5% high contamination risk, likely extend to the broader ecosystem.

7.3 Adoption by Category

Platform publishers (3 repos, ~50 skills): Anthropic (16 skills), OpenAI (32 skills for Codex), Vercel (5 skills). These serve as reference implementations.

Company-published (22+ repos, ~300+ skills): The largest segment. Microsoft alone publishes 143 skills covering Azure SDKs across five languages. Other major publishers include Sentry (16), HashiCorp (13), WordPress (13), Expo (9), Hugging Face (9), Cloudflare (9), and Vue.js (8). Companies publish skills to reduce integration friction for developers using their products.

Community collections (10+ repos, ~400+ skills): Multi-skill repositories from community maintainers. K-Dense-AI publishes 145+ scientific computing skills. Anthony Fu maintains 17 skills for the Vue/Nuxt/Vite ecosystem. Obsidian's CEO publishes 5 skills for the knowledge management tool. The K-Dense/Superpowers family covers development methodology.

Individual community skills (60+ repos, ~100+ skills): Single-purpose skills covering niche use cases: D3.js visualization, Playwright testing, ffuf security scanning, Home Assistant automation, video editing, and more.

Security-focused (5+ repos, ~65+ skills): Trail of Bits (52+ skills) dominates this space with vulnerability scanning, fuzzing, and audit skills. Prompt Security (7 skills) and Cisco AI Defense provide security scanning skills. Snyk publishes an agent security scanner.

Vertical/domain-specific (10+ repos, ~100+ skills):

- *Legal* (lawvable): 38 skills for contract review, compliance, and legal document drafting
- *Biotech* (Adaptyv Bio): 21 skills for protein design, AlphaFold, and computational biology
- *Embedded/IoT* (Zephyr): 21 skills for RTOS development, BLE, and board bringup
- *DevOps*: 6 skills for Terraform, Kubernetes, CI/CD, and monitoring
- *Business strategy* (wondelai): 25 skills covering Blue Ocean Strategy, Design Sprint, and similar frameworks

7.4 Ecosystem Infrastructure

Several supporting tools have emerged:

- **skills.sh**: Web registry with 58,000+ installations, providing search and discovery
- **Vercel `npx skills` CLI**: One-command installation of skills from any GitHub repository
- **SkillsMP**: Marketplace with quality indicators and category filtering

- **skill-validator** (Carey 2026): Structural validation tool (used in this analysis)
- **Cisco AI Defense skill-scanner**: Security scanner for detecting malicious skill patterns
- **Snyk agent-scan**: Security scanner for AI agents and skills

7.5 Implications for Our Findings

If the quality patterns we observe in our 673-skill sample hold across the full 1,400+ skill ecosystem:

- **~308 skills** may fail structural validation
- **~21 skills** may have structural indicators of high cross-contamination risk
- **~305 skills** may have structural indicators of medium contamination risk

These estimates show the need for quality standards and validation tooling. The ecosystem has reached a scale where manual review is impractical, and automated quality gates are necessary.

7.6 Source Repositories Not Included in Primary Analysis

For completeness, we list additional repositories identified during our survey that were not included in our primary analysis. These represent opportunities for future expansion:

- **OthmanAdi/planning-with-files** (13,888 stars): 13 planning workflow skills
- **blader/humanizer** (4,798 stars): AI text humanization
- **antfu/skills** (3,403 stars): Vue/Nuxt ecosystem (partially included)
- **blader/Claudeception** (1,624 stars): Autonomous skill extraction
- **CharlesWiltgen/Axiom** (465 stars): 144 Apple xOS development skills
- **daymade/claude-code-skills** (579 stars): 37 production-ready skills
- **Aaronontheweb/dotnet-skills** (327 stars): 33 .NET ecosystem skills
- **wondelai/skills** (104 stars): 25 business strategy skills
- **sundial-org/skills** (142 stars): 11 research-oriented skills
- Multiple awesome-lists curating 300-800+ skills each

References

"Agent Skills Specification." 2025. https://agentskills.io.

Ali, Ameen, Lior Wolf, and Ivan Titov. 2024. "Mitigating Copy Bias in in-Context Learning Through Neuron Pruning." *arXiv Preprint arXiv:2410.01288*.

An, Chenxin, Jun Zhang, Ming Zhong, et al. 2025. "Why Does the Effective Context Length of LLMs Fall Short?" *International Conference on Learning Representations (ICLR)*.

Anthropic. 2025. "Skill Creator: Guide for Creating Effective Skills." https://github.com/anthropics/skills/tree/main/skills/skill-creator.

Carey, Dachary. 2026. "Skill-Validator: A Go CLI for Validating Agent Skills." https://github.com/dacharyc/skill-validator.

Du, Yufeng, Minyang Tian, Srikanth Ronanki, et al. 2025. "Context Length Alone Hurts LLM Performance Despite Perfect Retrieval." *Findings of the Association for Computational Linguistics: EMNLP 2025*.

Hong, Kelly, Anton Troynikov, and Jeff Huber. 2025. *Context Rot: How Increasing Input Tokens Impacts LLM Performance*. Chroma. https://research.trychroma.com/context-rot.

Li, Jia, Ge Li, Chongyang Tao, et al. 2025. "Large Language Model-Aware in-Context Learning for Code Generation." *ACM Transactions on Software Engineering and Methodology*.

Liu, Nelson F., Kevin Lin, John Hewitt, et al. 2024. "Lost in the Middle: How Language Models Use Long Contexts." *Transactions of the Association for Computational Linguistics* 12: 157–73.

Liu, Yi, Weizhe Wang, Ruitao Feng, et al. 2026. "Agent Skills in the Wild: An Empirical Study of Security Vulnerabilities at Scale." *arXiv Preprint arXiv:2601.10338*.

Moumoula, Micheline Bénédicte, Serge Lionel Nikiema, Abdoul Kader Kaboré, Jacques Klein, and Tegawendé F. Bissyande. 2026. "Programming Language Confusion: When Code LLMs Can't Keep Their Languages Straight." *Proceedings of the IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*.

Salim, Mohamad, Jasmine Latendresse, SayedHassan Khatoonabadi, and Emad Shihab. 2026. "Tokenomics: Quantifying Where Tokens Are Used in Agentic Software Engineering." *Proceedings of the 23rd International Conference on Mining Software Repositories (MSR)*.

Shi, Freda, Xinyun Chen, Kanishka Misra, et al. 2023. "Large Language Models Can Be Easily Distracted by Irrelevant Context." *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 31210–27.

Tian, Yuan, and Tianyi Zhang. 2025. "Selective Prompt Anchoring for Code Generation." *International Conference on Machine Learning (ICML)*.

Xiao, Yuan-An, Pengfei Gao, Chao Peng, and Yingfei Xiong. 2025. "Improving the Efficiency of LLM Agent Systems Through Trajectory Reduction." *arXiv Preprint arXiv:2509.23586*.

Yoran, Ori, Tomer Wolfson, Ori Ram, and Jonathan Berant. 2024. "Making Retrieval-Augmented Language Models Robust to Irrelevant Context." *International Conference on Learning Representations (ICLR)*.